

Faster R-CNN Approach

Faster R-CNN

Import Libraries

Nedded Libraries

PyTorch

```
import torch
import torchvision
from torch.utils.data import Dataset, DataLoader
from torchvision import models
from torchvision.models.detection.faster_rcnn import
FastRCNNPredictor, fasterrcnn_resnet50_fpn
import albumentations as A
```

Image processing

```
from PIL import Image, ImageDraw, ExifTags, ImageColor, ImageFont
```

Image Plots

```
from matplotlib import pyplot as plt
import matplotlib.patches as patches
```

Data managements

```
import numpy as np
import pandas as pd
```

File interpretation

```
import os
import xml.etree.ElementTree as ET
import random
```

Others

```
import time
from collections import Counter
from random import seed, randint
from datetime import datetime
```

Find Annotation Files

Indicate the path for the annotation files.

Annotations directory path

```
ann_directory = '/content/FaceMaskDetection/annotations'
```

List directory

```
ann_files = os.listdir(ann_directory)
```

Find Image Files

Indicate the path for the image files.

```
# Image directory path
img_directory = '/content/FaceMaskDetection/images'
```

```
# List directory
img_files = os.listdir(img_directory)
```

Find Annotation Files

Indicate the path for the annotation files.

Helper Functions

This are auxiliary functions used throughout the notebook. This way the notebook stays tidy and clean.

```
def draw_bounding_boxes(img_tensor, target=None, prediction=None):
    """Draws bounding boxes in given images. Displays them

    Inputs:
        img:
            Image in tensor format.
        target:
            target dictionary containing bboxes list wit format ->
            [xmin, ymin, xmax, ymax]

    Returns:
        None
    """

    img = torchvision.transforms.ToPILImage()(img_tensor)

    # fetching the dimensions
    wid, hgt = img.size
    print(str(wid) + "x" + str(hgt))

    # Img to draw in
    draw = ImageDraw.Draw(img)

    if target:
        target_bboxes = target['boxes'].numpy().tolist()
        target_labels = decode_labels(target['labels'].numpy())

        for i in range(len(target_bboxes)):
            # Create Rectangle patches and add the patches to the axes
            draw.rectangle(target_bboxes[i], fill=None,
                outline='green', width=2)
```

```

        draw.text(target_bboxes[i][:2], target_labels[i],
fill='green', font=None, anchor=None, spacing=4,
                align='left', direction=None, features=None,
language=None, stroke_width=0, stroke_fill=None,
                embedded_color=False)

    if prediction:
        prediction_bboxes =
prediction['boxes'].detach().cpu().numpy().tolist()
        prediction_labels =
decode_labels(prediction['labels'].detach().cpu().numpy())
        for i in range(len(prediction_bboxes)):
            # Create Rectangle patches and add the patches to the axes
            draw.rectangle(prediction_bboxes[i], fill=None,
outline='red', width=2)
            draw.text(prediction_bboxes[i][:2], prediction_labels[i],
fill='red', font=None, anchor=None, spacing=4,
                    align='left', direction=None, features=None,
language=None, stroke_width=0, stroke_fill=None,
                    embedded_color=False)

```

```

display(img)

```

```

def encoded_labels(lst_labels):
    """Encodes label classes from string to integers.

    Labels are encoded accordingly:
        - background => 0
        - with_mask => 1
        - mask_weared_incorrect => 2
        - without_mask => 3

    Args:
        lst_labels:
            A list with classes in string format (e.g.
['with_mask', 'mask_weared_incorrect'...]).

    Returns:
        encoded:
            A list with integers that represent each class.
    """

```

```

encoded=[]
for label in lst_labels:
    if label == "with_mask":
        code = 1
    elif label == "mask_weared_incorrect":
        code = 2
    elif label == "without_mask":
        code = 3

```

```

        else:
            code = 0
            encoded.append(code)
    return encoded

def decode_labels(lst_labels):
    """
    Decode label classes from integers to strings.
    Labels are encoded accordingly:
        - background => 0
        - with_mask => 1
        - mask_wearred_incorrect => 2
        - without_mask => 3

    Args:
        lst_labels:
            A list with classes in integer format (e.g. [1, 2, ...]).

    Returns:
        A list with strings that represent each class.
    """

    labels=[]
    for code in lst_labels:
        if code == 1:
            label = "with_mask"
        elif code == 2:
            label = "mask_wearred_incorrect"
        elif code == 3:
            label = "without_mask"
        else:
            label = 'background'
        labels.append(label)
    return labels

def build_model(nclasses):
    """
    Builds model. Uses Faster R-CNN pre-trained on COCO dataset.

    Args:
        nclasses:
            number of classes

    Return:
        model: Faster R-CNN pre-trained model
    """
    # load pre-trained model on COCO
    model = fasterrcnn_resnet50_fpn(pretrained=True, min_size=400,
max_size=700)

```

```

    # get the number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features

    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
nclasses)

    return model

def train_model(model, loader, optimizer, scheduler, epochs, device):
    """
    Inputs:
        - model
        - loader: Dataloader PyTorch object with training data
        - optimizer
        - scheduler
        - epochs
        - device

    Returns:
        - model
        - loss_list: list with mean loss per epoch. Epoch 1 is in index
0.
    """
    # Create a loss list to keep epoch average loss
    loss_list = []
    # Epochs
    for epoch in range(epochs):
        print('Starting epoch..... {}/{}'.format(epoch + 1, epochs))
        iteration = 0
        loss_sub_list = []
        start = time.time()
        for images, targets in loader:
            # Agregate images in batch loader
            images = list(image.to(device) for image in images)

            # Agregate targets in batch loader
            targets = [{key: val.to(device) for key, val in
target.items()} for target in targets]

            # Sets model to train mode (just a flag)
            model.train()

            # Output of model returns loss and detections
            optimizer.zero_grad()
            output = model(images, targets)

            # Calculate Cost
            losses = sum(loss for loss in output.values())
            loss_value = losses.item()

```

```

        loss_sub_list.append(loss_value)
        print('')

        # Update optimizer and learning rate
        losses.backward()
        optimizer.step()
        iteration += 1
        print('Iteration: {:d} --> Loss: {:.3f}'.format(iteration,
loss_value))

    end = time.time()
    # update scheduler
    scheduler.step()
    # print the loss of epoch
    epoch_loss = np.mean(loss_sub_list)
    loss_list.append(epoch_loss)
    print('Epoch loss: {:.3f} , time used:
({:.1f}s)'.format(epoch_loss, end - start))

    return model, loss_list

def apply_nms(orig_prediction, iou_thresh):
    """
    Applies non max supression and eliminates low score bounding
boxes.

    Args:
        orig_prediction: the model output. A dictionary containing
element scores and boxes.
        iou_thresh: Intersection over Union threshold. Every bbox
prediction with an IoU greater than this value
gets deleted in MMS.

    Returns:
        final_prediction: Resulting prediction
    """

    # torchvision returns the indices of the bboxes to keep
    keep = torchvision.ops.nms(orig_prediction['boxes'],
orig_prediction['scores'], iou_thresh)

    # Keep indices from nms
    final_prediction = orig_prediction
    final_prediction['boxes'] = final_prediction['boxes'][keep]
    final_prediction['scores'] = final_prediction['scores'][keep]
    final_prediction['labels'] = final_prediction['labels'][keep]

    return final_prediction

```

```

def remove_low_score_bb(orig_prediction, score_thresh):
    """
    Eliminates low score bounding boxes.

    Args:
        orig_prediction: the model output. A dictionary containing
        element scores and boxes.
        score_thresh: Boxes with a lower confidence score than this
        value get deleted

    Returns:
        final_prediction: Resulting prediction
    """

    # Remove low confidence scores according to given threshold
    index_list_scores = []
    scores = orig_prediction['scores'].detach().cpu().numpy()
    for i in range(len(scores)):
        if scores[i] > score_thresh:
            index_list_scores.append(i)
    keep = torch.tensor(index_list_scores)

    # Keep indices from high score bb
    final_prediction = orig_prediction
    final_prediction['boxes'] = final_prediction['boxes'][keep]
    final_prediction['scores'] = final_prediction['scores'][keep]
    final_prediction['labels'] = final_prediction['labels'][keep]

    return final_prediction

def collate_fn(batch):
    # Collate function for Dataloader
    return tuple(zip(*batch))

def IOU(box1, box2):
    """
    Intersection over Union - IoU

    *-----
    |      (x2min,y2min)
    |      *-----
    |      | ##### |
    |-----* (x1max,y1max)
    |
    |-----
    """

    Args:
        box1: [xmin,ymin,xmax,ymax]
        box2: [xmin,ymin,xmax,ymax]

    Returns:

```

```

    iou -> value of intersection over union of the 2 boxes
'''

# Compute coordinates of intersection
xmin_inter = max(box1[0], box2[0])
ymin_inter = max(box1[1], box2[1])
xmax_inter = min(box1[2], box2[2])
ymax_inter = min(box1[3], box2[3])

# calculate area of intersection rectangle
inter_area = max(0, xmax_inter - xmin_inter + 1) * max(0,
ymax_inter - ymin_inter + 1) # FIXME why plus one?

# calculate boxes areas
area1 = (box1[2] - box1[0] + 1) * (box1[3] - box1[1] + 1)
area2 = (box2[2] - box2[0] + 1) * (box2[3] - box2[1] + 1)

# compute IoU
iou = inter_area / float(area1 + area2 - inter_area)
assert iou >= 0
return iou

def compute_AP(ground_truth, predictions, iou_thresh=0.5,
n_classes=4):
    """
    Calculates Average Precision across all classes.

    Args:
        ground_truth: list with ground-truth objects. Needs to have
the following format: [sequence, frame, obj, [xmin, ymin, xmax, ymax],
label, score]
        predictions: list with predictions objects. Needs to have the
following format: [sequence, frame, obj, [xmin, ymin, xmax, ymax],
label, score]
        iou_thresh: IoU to which a prediction compared to a ground-
truth is considered right.
        n_classes: number of existent classes

    Returns:
        Average precision for the specified threshold.
    """
    # Initialize lists
    APs = []
    class_gt = []
    class_predictions = []

    # AP is computed for each class
    for c in range(n_classes):
        # Find gt and predictions of the class

```



```

for gt in ground_truth:
    if gt[4] == c:
        class_gt.append(gt)
for predict in predictions:
    if predict[4] == c:
        class_predictions.append(predict)

# Create dict with array of zeros for bb in each image
gt_amount_bb = Counter([gt[1] for gt in class_gt])
for key, val in gt_amount_bb.items():
    gt_amount_bb[key] = np.zeros(val)

# Sort class predictions by their score
class_predictions = sorted(class_predictions, key=lambda x:
x[5], reverse=True)

# Create arrays for Positives (True and False)
TP = np.zeros(len(class_predictions))
FP = np.zeros(len(class_predictions))
# Number of true boxes
truth = len(class_gt)

# Initializing aux variables
epsilon = 1e-6

# Iterate over predictions in each image and compare with
ground truth
for predict_idx, prediction in enumerate(class_predictions):
    # Filter prediction image ground truths
    image_gt = [obj for obj in class_gt if obj[1] ==
prediction[1]]

    # Initializing aux variables
    best_iou = -1
    best_gt_iou_idx = -1

    # Iterate through image ground truths and calculate IoUs
    for gt_idx, gt in enumerate(image_gt):
        iou = IOU(prediction[3], gt[3])
        if iou > best_iou:
            best_iou = iou
            best_gt_iou_idx = gt_idx

    # If the best IoU is greater than thresh than an TP
prediction has been found
    if best_iou > iou_thresh and best_gt_iou_idx > -1:
        # Check if gt box was already covered
        if gt_amount_bb[prediction[1]][best_gt_iou_idx] == 0:
            gt_amount_bb[prediction[1]][best_gt_iou_idx] = 1

```

```

# set as covered
        TP[predict_idx] = 1 # Count as true positive
    else:
        FP[predict_idx] = 1
    else:
        FP[predict_idx] = 1

    # Calculate recall and precision
    TP_cumsum = np.cumsum(TP)
    FP_cumsum = np.cumsum(FP)
    recall = np.append([0], TP_cumsum / (truth + epsilon))
    precision = np.append([1], np.divide(TP_cumsum, (TP_cumsum +
FP_cumsum + epsilon)))

    # Calculate the area precision/recall and add to list
    APs.append(np.trapz(precision, recall))

    return sum(APs)/len(APs) # average of class precisions

def compute_mAP(ground_truth, predictions, n_classes):
    """
    Calls AP computation for different levels of IoUs, [0.5:.05:0.95].

    Args:
        ground_truth: list with ground-truth objects. Needs to have
the following format: [sequence, frame, obj, [xmin, ymin, xmax, ymax],
label, score]
        predictions: list with predictions objects. Needs to have the
following format: [sequence, frame, obj, [xmin, ymin, xmax, ymax],
label, score]
        n_classes: number of existent classes.

    Returns:
        mAp and list with APs for each IoU threshold.
    """
    # return mAP
    APs = [compute_AP(ground_truth, predictions, iou_thresh,
n_classes) for iou_thresh in np.arange(0.5, 1.0, 0.05)]
    return np.mean(APs), APs

@torch.no_grad()
def evaluate(model, data_loader, device, sequences=1):
    """
    Evaluates model mAP for IoU range of [0.5:.05:0.95].

    Args:
        model: -
        data_loader: -
        device: -

```

sequences: the number of sequences of images to pass, if any

Returns:

mAP and AP list for each IoU threshold in range [0.5:.05:0.95]
"""

```
# Set evaluation mode flag
model.eval()
# Create list with all object detection -> [set, frame, obj,
[xmin,ymin,xmax,ymax], label, score]
ground_truth = []
predictions = []

# Gather all targets and outputs on test set
for image, targets in data_loader:
    image = [img.to(device) for img in image]
    outputs = model(image)
    for idx in range(len(outputs)):
        outputs[idx] = apply_nms(outputs[idx], iou_thresh=0.5)

# create list for targets and outputs to pass to compute_mAP()
# lists have the following structure: [sequence, frame,
obj_idx, [xmin, ymin, xmax, ymax], label, score]
for s in range(sequences):
    obj_gt = 0
    obj_target = 0
    for out, target in zip(outputs, targets):

        for i in range(len(target['boxes'])):
            ground_truth.append([s,
target['image_id'].detach().cpu().numpy()[0], obj_target,

target['boxes'].detach().cpu().numpy()[i],

target['labels'].detach().cpu().numpy()[i], 1])
            obj_target += 1

        for j in range(len(out['boxes'])):
            predictions.append([s,
target['image_id'].detach().cpu().numpy()[0], obj_gt,

out['boxes'].detach().cpu().numpy()[j],

out['labels'].detach().cpu().numpy()[j],

out['scores'].detach().cpu().numpy()[j]])
            obj_gt += 1

mAP, AP = compute_mAP(ground_truth, predictions, n_classes=4)
```

```

print("mAP:{:.3f}".format(mAP))
for ap_metric, iou in zip(AP, np.arange(0.5, 1, 0.05)):
    print("\tAP at IoU level [{:.2f}]: {:.3f}".format(iou,
ap_metric))

return mAP, AP

```

Create Dataset Class

Dataset class to feed the dataloader.

```

# Create dataset object
class MyDataset(Dataset):

    # Constructor
    def __init__(self, ann_dir, img_dir, transform=None,
mode='train'):

        # Image directories
        self.ann_dir = ann_dir
        self.img_dir = img_dir

        # The transform is going to be used on image
        self.transform = transform

        # Create dataframe to hold info
        self.data = pd.DataFrame(columns=['Filename', 'BoundingBoxes',
'Labels', 'Area', 'N_Objects'])

        # Append rows with image filename and respective bounding
boxes to the df
        for file in enumerate(os.listdir(img_dir)):

            # Find image annotation file
            ann_file_path = os.path.join(ann_dir, file[1][:-4]) +
'.xml'

            # Read XML file and return bounding boxes and class
attributes
            objects = self.read_XML_classf(ann_file_path)

            # Create list of labels in an image
            list_labels = encoded_labels(objects[0]['labels'])

            # Create list of bounding boxes in an image
            list_bb = []
            list_area = []
            n_obj = len(objects[0]['objects'])
            for i in objects[0]['objects']:
                list = [i['xmin'], i['ymin'], i['xmax'], i['ymax']]

```

```

        list_bb.append(list)
        list_area.append((i['xmax'] - i['xmin']) * (i['ymax']
- i['ymin']))

        # Create dataframe object with row containing [(Image file
name),(Bounding Box List)]
        df = pd.DataFrame([[file[1], list_bb, list_labels,
list_area, n_obj]],
                           columns=['Filename', 'BoundingBoxes',
'Labels', 'Area', 'N_Objects'])
        self.data = self.data.append(df)

        if mode == 'train':
            self.data = self.data[:680]
        elif mode == 'validation':
            self.data = self.data[680:700]
        elif mode == 'test':
            self.data = self.data[700:850]

        # Number of images in dataset
        self.len = self.data.shape[0]

        # Get the length

def __len__(self):
    return self.len

# Getter
def __getitem__(self, idx):

    # Image file path
    img_name = os.path.join(self.img_dir, self.data.iloc[idx, 0])

    # Open image file and tranform to tensor
    img = Image.open(img_name).convert('RGB')

    # Get bounding box coordinates
    bbox = torch.tensor(self.data.iloc[idx, 1])

    # Get labels
    labels = torch.tensor(self.data.iloc[idx, 2])

    # Get bounding box areas
    area = torch.tensor(self.data.iloc[idx, 3])

    # If any, aplly tranformations to image and bounding box mask
    if self.transform:
        # Convert PIL image to numpy array
        img = np.array(img)

```

```

        # Apply transformations
        transformed = self.transform(image=img, bboxes=bbbox)
        # Convert numpy array to PIL Image
        img = Image.fromarray(transformed['image'])
        # Get transformed bb
        bbbox = torch.tensor(transformed['bboxes'])

    # suppose all instances are not crowd
    num_objs = self.data.iloc[idx, 4]
    iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

    # Transform img to tensor
    img = torchvision.transforms.ToTensor()(img)

    # Build Targer dict
    target= {"boxes": bbbox, "labels": labels, "image_id":
torch.tensor([idx]), "area": area, "iscrowd": iscrowd}

    return img, target

# XML reader -> returns dictionary with image bounding boxes sizes
def read_XML_classf(self, ann_file_path):
    bboxes = [{
        'file': ann_file_path,
        'labels': [],
        'objects': []
    }]

    # Reading XML file objects and print Bounding Boxes
    tree = ET.parse(ann_file_path)
    root = tree.getroot()
    objects = root.findall('object')

    for obj in objects:
        # label
        label = obj.find('name').text
        bboxes[0]['labels'].append(label)

        # bbox dimensions
        bndbox = obj.find('bndbox')
        xmin = int(bndbox.find('xmin').text)
        ymin = int(bndbox.find('ymin').text)
        xmax = int(bndbox.find('xmax').text)
        ymax = int(bndbox.find('ymax').text)
        bboxes[0]['objects'].append({'xmin': xmin, 'ymin': ymin,
'xmax': xmax, 'ymax': ymax})

    return bboxes

```

Create Data Pipeline

Create Data Pipeline

Training Data

```
dataset_train = MyDataset(ann_directory,img_directory, mode = 'train')
loader_train = DataLoader(dataset_train, batch_size=4, shuffle=True,
collate_fn=collate_fn)
```

Validation Data

```
dataset_validation = MyDataset(ann_directory,img_directory, mode =
'validation')
loader_val = DataLoader(dataset_validation, batch_size=4,
shuffle=True, collate_fn=collate_fn)
```

Test Data

```
dataset_test = MyDataset(ann_directory,img_directory, mode = 'test')
loader_test = DataLoader(dataset_test, batch_size=4, shuffle=True,
collate_fn=collate_fn)
```

Test if dataset is working correctly. Print out ground truth bounding box of first image.

pick one image from the train set

```
img, target = dataset_train[0]
draw_bounding_boxes(img, target)
```


```
img, target = dataset_train[4]
draw_bounding_boxes(img, target)
```

```
img, target = dataset_train[7]
draw_bounding_boxes(img, target)
```

400x210



301x400



267x400



Setting up the Faster R-CNN Model

Setting up GPU device

```
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')
```

N° of classes: background, with_mask, mask_wearred_incorrect, without_mask and build model (faster r-cnn)

```
num_classes = 4
```

```
model = build_model(num_classes)
```

```
model = model.to(device)
```

```
/usr/local/lib/python3.7/dist-packages/torchvision/models/_
```

```
utils.py:209: UserWarning: The parameter 'pretrained' is deprecated
since 0.13 and will be removed in 0.15, please use 'weights' instead.
```

```
    f"The parameter '{pretrained_param}' is deprecated since 0.13 and
will be removed in 0.15, "
```

```
/usr/local/lib/python3.7/dist-packages/torchvision/models/_utils.py:22
```

3: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and will be removed in 0.15. The current behavior is equivalent to passing `weights=FasterRCNN_ResNet50_FPN_Weights.COCO_V1`. You can also use `weights=FasterRCNN_ResNet50_FPN_Weights.DEFAULT` to get the most up-to-date weights.

warnings.warn(msg)

Downloading:

"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth" to
/root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth

```
{"version_major":2,"version_minor":0,"model_id":"5d22a8e2c3dd4234b40ea4f6e7a9dc3a"}
```

Set Hyper-parameters

Network params

```
params = [p for p in model.parameters() if p.requires_grad]
```

Optimizers

```
optimizer = torch.optim.Adam(params, lr=0.0001)
```

```
#optimizer = torch.optim.SGD(params, lr=0.005)
```

Learning Rate, lr decreases to half every 2 epochs

```
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=2,  
gamma=0.5)
```

Number of epochs to perform

```
epochs=20
```

Train the model

```
pip install
```

```
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-  
master#subdirectory=PythonAPI
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting

```
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-  
master#subdirectory=PythonAPI
```

Cloning <https://github.com/gautamchitnis/cocoapi.git> (to revision cocodataset-master) to /tmp/pip-req-build-wzjl5jf

Running command git clone -q

```
https://github.com/gautamchitnis/cocoapi.git /tmp/pip-req-build-  
wzjl5jf
```

Running command git checkout -b cocodataset-master --track origin/cocodataset-master

Switched to a new branch 'cocodataset-master'

```
Branch 'cocodataset-master' set up to track remote branch
'cocodataset-master' from 'origin'.
Building wheels for collected packages: pycocotools
  Building wheel for pycocotools (setup.py) ... e=pycocotools-2.0-
cp37-cp37m-linux_x86_64.whl size=265320
sha256=3abale2d4e84d0edca62fc1d4cf960c07e7dea2d21a5d64744e0b3e85b3a979
6
  Stored in directory:
/tmp/pip-ephem-wheel-cache-vs71d2py/wheels/6e/c9/59/56484d4d5ac1ab292a
452b4c3870277256551505954fc4a1db
Successfully built pycocotools
Installing collected packages: pycocotools
  Attempting uninstall: pycocotools
    Found existing installation: pycocotools 2.0.4
    Uninstalling pycocotools-2.0.4:
      Successfully uninstalled pycocotools-2.0.4
Successfully installed pycocotools-2.0
```

```
!git clone https://github.com/pytorch/vision.git
%cd vision
!git checkout v0.3.0
```

```
!cp references/detection/utils.py ../
!cp references/detection/transforms.py ../
!cp references/detection/coco_eval.py ../
!cp references/detection/engine.py ../
!cp references/detection/coco_utils.py ../
%cd ..
```

```
Cloning into 'vision'...
remote: Enumerating objects: 190176, done.ote: Counting objects: 100%
(33/33), done.ote: Compressing objects: 100% (26/26), done.ote: Total
190176 (delta 11), reused 12 (delta 7), pack-reused 190143ake
experimental
changes and commit them, and you can discard any commits you make in
this
state without impacting any branches by performing another checkout.
```

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-b` with the checkout command again. Example:

```
git checkout -b <new-branch-name>
```

HEAD is now at be376084d version check against PyTorch's CUDA version /content

```
from engine import train_one_epoch
# Training
```

```

for epoch in range(epochs):
    # train for one epoch, printing every 50 iterations
    train_one_epoch(model, optimizer, loader_train, device, epoch,
print_freq=20)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, loader_val, device=device)

Epoch: [0] [ 0/170] eta: 0:02:57 lr: 0.000001 loss: 2.1042
(2.1042) loss_classifier: 1.6616 (1.6616) loss_box_reg: 0.0811
(0.0811) loss_objectness: 0.2105 (0.2105) loss_rpn_box_reg: 0.1509
(0.1509) time: 1.0424 data: 0.1549 max mem: 1391
Epoch: [0] [ 20/170] eta: 0:01:06 lr: 0.000013 loss: 1.7572
(1.6751) loss_classifier: 1.2356 (1.2573) loss_box_reg: 0.1703
(0.1726) loss_objectness: 0.1480 (0.1799) loss_rpn_box_reg: 0.0220
(0.0654) time: 0.4127 data: 0.1111 max mem: 2507
Epoch: [0] [ 40/170] eta: 0:00:55 lr: 0.000024 loss: 0.7379
(1.2276) loss_classifier: 0.3061 (0.8143) loss_box_reg: 0.3008
(0.2367) loss_objectness: 0.0648 (0.1258) loss_rpn_box_reg: 0.0233
(0.0508) time: 0.4027 data: 0.1004 max mem: 2507
Epoch: [0] [ 60/170] eta: 0:00:45 lr: 0.000036 loss: 0.4773
(1.0015) loss_classifier: 0.1737 (0.6075) loss_box_reg: 0.2569
(0.2530) loss_objectness: 0.0275 (0.0975) loss_rpn_box_reg: 0.0105
(0.0435) time: 0.4028 data: 0.1013 max mem: 2507
Epoch: [0] [ 80/170] eta: 0:00:37 lr: 0.000048 loss: 0.4473
(0.8779) loss_classifier: 0.1340 (0.4972) loss_box_reg: 0.2861
(0.2585) loss_objectness: 0.0280 (0.0830) loss_rpn_box_reg: 0.0167
(0.0391) time: 0.4163 data: 0.1110 max mem: 2697
Epoch: [0] [100/170] eta: 0:00:28 lr: 0.000060 loss: 0.3520
(0.7831) loss_classifier: 0.1166 (0.4248) loss_box_reg: 0.2058
(0.2511) loss_objectness: 0.0146 (0.0723) loss_rpn_box_reg: 0.0065
(0.0348) time: 0.4034 data: 0.1049 max mem: 2697
Epoch: [0] [120/170] eta: 0:00:20 lr: 0.000072 loss: 0.4257
(0.7362) loss_classifier: 0.1289 (0.3801) loss_box_reg: 0.2208
(0.2515) loss_objectness: 0.0322 (0.0687) loss_rpn_box_reg: 0.0236
(0.0359) time: 0.4051 data: 0.1056 max mem: 2697
Epoch: [0] [140/170] eta: 0:00:12 lr: 0.000083 loss: 0.4665
(0.6990) loss_classifier: 0.1249 (0.3479) loss_box_reg: 0.2524
(0.2508) loss_objectness: 0.0144 (0.0653) loss_rpn_box_reg: 0.0062
(0.0349) time: 0.4090 data: 0.1078 max mem: 2697
Epoch: [0] [160/170] eta: 0:00:04 lr: 0.000095 loss: 0.3958
(0.6606) loss_classifier: 0.1081 (0.3204) loss_box_reg: 0.2100
(0.2466) loss_objectness: 0.0151 (0.0599) loss_rpn_box_reg: 0.0079
(0.0337) time: 0.4144 data: 0.1150 max mem: 2697
Epoch: [0] [169/170] eta: 0:00:00 lr: 0.000100 loss: 0.3870
(0.6412) loss_classifier: 0.1127 (0.3085) loss_box_reg: 0.2100
(0.2428) loss_objectness: 0.0113 (0.0573) loss_rpn_box_reg: 0.0127
(0.0326) time: 0.4277 data: 0.1202 max mem: 2794
Epoch: [0] Total time: 0:01:10 (0.4141 s / it)

```

mAP:0.390

AP at IoU level [0.50]: 0.666
AP at IoU level [0.55]: 0.666
AP at IoU level [0.60]: 0.632
AP at IoU level [0.65]: 0.608
AP at IoU level [0.70]: 0.549
AP at IoU level [0.75]: 0.479
AP at IoU level [0.80]: 0.226
AP at IoU level [0.85]: 0.060
AP at IoU level [0.90]: 0.017
AP at IoU level [0.95]: 0.000

Epoch: [1] [0/170] eta: 0:00:57 lr: 0.000100 loss: 0.4972
(0.4972) loss_classifier: 0.1732 (0.1732) loss_box_reg: 0.2865
(0.2865) loss_objectness: 0.0262 (0.0262) loss_rpn_box_reg: 0.0114
(0.0114) time: 0.3408 data: 0.0574 max mem: 2794
Epoch: [1] [20/170] eta: 0:00:54 lr: 0.000100 loss: 0.2797
(0.2914) loss_classifier: 0.0942 (0.0968) loss_box_reg: 0.1649
(0.1723) loss_objectness: 0.0071 (0.0112) loss_rpn_box_reg: 0.0068
(0.0112) time: 0.3628 data: 0.0580 max mem: 2794
Epoch: [1] [40/170] eta: 0:00:47 lr: 0.000100 loss: 0.3508
(0.3219) loss_classifier: 0.0933 (0.1016) loss_box_reg: 0.2163
(0.1905) loss_objectness: 0.0141 (0.0132) loss_rpn_box_reg: 0.0113
(0.0165) time: 0.3645 data: 0.0596 max mem: 2794
Epoch: [1] [60/170] eta: 0:00:39 lr: 0.000100 loss: 0.3592
(0.3315) loss_classifier: 0.1031 (0.1014) loss_box_reg: 0.2025
(0.1971) loss_objectness: 0.0129 (0.0136) loss_rpn_box_reg: 0.0137
(0.0194) time: 0.3613 data: 0.0630 max mem: 2794
Epoch: [1] [80/170] eta: 0:00:32 lr: 0.000100 loss: 0.3977
(0.3441) loss_classifier: 0.1186 (0.1052) loss_box_reg: 0.2242
(0.2050) loss_objectness: 0.0078 (0.0141) loss_rpn_box_reg: 0.0084
(0.0199) time: 0.3697 data: 0.0637 max mem: 2794
Epoch: [1] [100/170] eta: 0:00:25 lr: 0.000100 loss: 0.3831
(0.3510) loss_classifier: 0.1259 (0.1078) loss_box_reg: 0.2064
(0.2076) loss_objectness: 0.0197 (0.0148) loss_rpn_box_reg: 0.0222
(0.0208) time: 0.3629 data: 0.0638 max mem: 2794
Epoch: [1] [120/170] eta: 0:00:18 lr: 0.000100 loss: 0.2674
(0.3452) loss_classifier: 0.0875 (0.1064) loss_box_reg: 0.1748
(0.2031) loss_objectness: 0.0060 (0.0150) loss_rpn_box_reg: 0.0063
(0.0207) time: 0.3721 data: 0.0601 max mem: 2794
Epoch: [1] [140/170] eta: 0:00:10 lr: 0.000100 loss: 0.2462
(0.3347) loss_classifier: 0.0754 (0.1033) loss_box_reg: 0.1627
(0.1971) loss_objectness: 0.0050 (0.0143) loss_rpn_box_reg: 0.0042
(0.0200) time: 0.3575 data: 0.0575 max mem: 2794
Epoch: [1] [160/170] eta: 0:00:03 lr: 0.000100 loss: 0.3389
(0.3334) loss_classifier: 0.1002 (0.1035) loss_box_reg: 0.1809
(0.1963) loss_objectness: 0.0060 (0.0140) loss_rpn_box_reg: 0.0072
(0.0197) time: 0.3634 data: 0.0663 max mem: 2794
Epoch: [1] [169/170] eta: 0:00:00 lr: 0.000100 loss: 0.2710
(0.3316) loss_classifier: 0.0872 (0.1024) loss_box_reg: 0.1671
(0.1956) loss_objectness: 0.0082 (0.0140) loss_rpn_box_reg: 0.0108

(0.0196) time: 0.3661 data: 0.0653 max mem: 2794

Epoch: [1] Total time: 0:01:01 (0.3642 s / it)

mAP:0.426

AP at IoU level [0.50]: 0.698

AP at IoU level [0.55]: 0.675

AP at IoU level [0.60]: 0.654

AP at IoU level [0.65]: 0.642

AP at IoU level [0.70]: 0.591

AP at IoU level [0.75]: 0.532

AP at IoU level [0.80]: 0.346

AP at IoU level [0.85]: 0.120

AP at IoU level [0.90]: 0.004

AP at IoU level [0.95]: 0.000

Epoch: [2] [0/170] eta: 0:01:03 lr: 0.000050 loss: 0.4140

(0.4140) loss_classifier: 0.1041 (0.1041) loss_box_reg: 0.2435

(0.2435) loss_objectness: 0.0194 (0.0194) loss_rpn_box_reg: 0.0470

(0.0470) time: 0.3708 data: 0.0549 max mem: 2794

Epoch: [2] [20/170] eta: 0:00:53 lr: 0.000050 loss: 0.1858

(0.2431) loss_classifier: 0.0538 (0.0775) loss_box_reg: 0.1200

(0.1438) loss_objectness: 0.0046 (0.0092) loss_rpn_box_reg: 0.0042

(0.0127) time: 0.3556 data: 0.0624 max mem: 2794

Epoch: [2] [40/170] eta: 0:00:46 lr: 0.000050 loss: 0.2768

(0.2706) loss_classifier: 0.0788 (0.0814) loss_box_reg: 0.1706

(0.1631) loss_objectness: 0.0050 (0.0100) loss_rpn_box_reg: 0.0082

(0.0161) time: 0.3651 data: 0.0626 max mem: 2794

Epoch: [2] [60/170] eta: 0:00:40 lr: 0.000050 loss: 0.2187

(0.2577) loss_classifier: 0.0596 (0.0795) loss_box_reg: 0.1343

(0.1554) loss_objectness: 0.0046 (0.0087) loss_rpn_box_reg: 0.0044

(0.0141) time: 0.3704 data: 0.0620 max mem: 2794

Epoch: [2] [80/170] eta: 0:00:32 lr: 0.000050 loss: 0.2351

(0.2553) loss_classifier: 0.0701 (0.0788) loss_box_reg: 0.1365

(0.1529) loss_objectness: 0.0030 (0.0077) loss_rpn_box_reg: 0.0058

(0.0158) time: 0.3673 data: 0.0662 max mem: 2794

Epoch: [2] [100/170] eta: 0:00:25 lr: 0.000050 loss: 0.2451

(0.2517) loss_classifier: 0.0620 (0.0765) loss_box_reg: 0.1431

(0.1520) loss_objectness: 0.0034 (0.0075) loss_rpn_box_reg: 0.0099

(0.0157) time: 0.3640 data: 0.0623 max mem: 2794

Epoch: [2] [120/170] eta: 0:00:18 lr: 0.000050 loss: 0.1962

(0.2487) loss_classifier: 0.0550 (0.0754) loss_box_reg: 0.1335

(0.1510) loss_objectness: 0.0021 (0.0073) loss_rpn_box_reg: 0.0043

(0.0150) time: 0.3581 data: 0.0644 max mem: 2794

Epoch: [2] [140/170] eta: 0:00:10 lr: 0.000050 loss: 0.2371

(0.2492) loss_classifier: 0.0730 (0.0751) loss_box_reg: 0.1473

(0.1514) loss_objectness: 0.0042 (0.0073) loss_rpn_box_reg: 0.0055

(0.0154) time: 0.3608 data: 0.0573 max mem: 2794

Epoch: [2] [160/170] eta: 0:00:03 lr: 0.000050 loss: 0.1850

(0.2435) loss_classifier: 0.0500 (0.0733) loss_box_reg: 0.1216

(0.1482) loss_objectness: 0.0036 (0.0070) loss_rpn_box_reg: 0.0041

(0.0150) time: 0.3638 data: 0.0655 max mem: 2794

Epoch: [2] [169/170] eta: 0:00:00 lr: 0.000050 loss: 0.2194

(0.2432) loss_classifier: 0.0689 (0.0742) loss_box_reg: 0.1316
(0.1474) loss_objectness: 0.0036 (0.0070) loss_rpn_box_reg: 0.0036
(0.0146) time: 0.3621 data: 0.0620 max mem: 2794
Epoch: [2] Total time: 0:01:01 (0.3632 s / it)
mAP:0.455
AP at IoU level [0.50]: 0.710
AP at IoU level [0.55]: 0.694
AP at IoU level [0.60]: 0.680
AP at IoU level [0.65]: 0.654
AP at IoU level [0.70]: 0.634
AP at IoU level [0.75]: 0.579
AP at IoU level [0.80]: 0.428
AP at IoU level [0.85]: 0.151
AP at IoU level [0.90]: 0.020
AP at IoU level [0.95]: 0.000
Epoch: [3] [0/170] eta: 0:01:05 lr: 0.000050 loss: 0.3210
(0.3210) loss_classifier: 0.0884 (0.0884) loss_box_reg: 0.1770
(0.1770) loss_objectness: 0.0354 (0.0354) loss_rpn_box_reg: 0.0203
(0.0203) time: 0.3869 data: 0.0680 max mem: 2794
Epoch: [3] [20/170] eta: 0:00:54 lr: 0.000050 loss: 0.1955
(0.2044) loss_classifier: 0.0585 (0.0639) loss_box_reg: 0.1187
(0.1210) loss_objectness: 0.0039 (0.0080) loss_rpn_box_reg: 0.0046
(0.0116) time: 0.3602 data: 0.0600 max mem: 2794
Epoch: [3] [40/170] eta: 0:00:47 lr: 0.000050 loss: 0.1854
(0.2003) loss_classifier: 0.0481 (0.0625) loss_box_reg: 0.1341
(0.1216) loss_objectness: 0.0014 (0.0059) loss_rpn_box_reg: 0.0035
(0.0102) time: 0.3660 data: 0.0597 max mem: 2794
Epoch: [3] [60/170] eta: 0:00:39 lr: 0.000050 loss: 0.2053
(0.2137) loss_classifier: 0.0722 (0.0670) loss_box_reg: 0.1224
(0.1277) loss_objectness: 0.0025 (0.0064) loss_rpn_box_reg: 0.0050
(0.0127) time: 0.3599 data: 0.0635 max mem: 2794
Epoch: [3] [80/170] eta: 0:00:32 lr: 0.000050 loss: 0.2003
(0.2166) loss_classifier: 0.0667 (0.0670) loss_box_reg: 0.1350
(0.1294) loss_objectness: 0.0009 (0.0057) loss_rpn_box_reg: 0.0067
(0.0144) time: 0.3748 data: 0.0629 max mem: 2794
Epoch: [3] [100/170] eta: 0:00:25 lr: 0.000050 loss: 0.1670
(0.2123) loss_classifier: 0.0511 (0.0647) loss_box_reg: 0.1104
(0.1283) loss_objectness: 0.0019 (0.0059) loss_rpn_box_reg: 0.0030
(0.0133) time: 0.3570 data: 0.0602 max mem: 2794
Epoch: [3] [120/170] eta: 0:00:18 lr: 0.000050 loss: 0.1593
(0.2078) loss_classifier: 0.0503 (0.0631) loss_box_reg: 0.1078
(0.1265) loss_objectness: 0.0023 (0.0054) loss_rpn_box_reg: 0.0030
(0.0128) time: 0.3671 data: 0.0584 max mem: 2794
Epoch: [3] [140/170] eta: 0:00:10 lr: 0.000050 loss: 0.2051
(0.2096) loss_classifier: 0.0630 (0.0632) loss_box_reg: 0.1269
(0.1278) loss_objectness: 0.0051 (0.0057) loss_rpn_box_reg: 0.0102
(0.0129) time: 0.3494 data: 0.0605 max mem: 2794
Epoch: [3] [160/170] eta: 0:00:03 lr: 0.000050 loss: 0.1918
(0.2091) loss_classifier: 0.0472 (0.0626) loss_box_reg: 0.1245
(0.1280) loss_objectness: 0.0037 (0.0056) loss_rpn_box_reg: 0.0034

(0.0129) time: 0.3685 data: 0.0652 max mem: 2794
Epoch: [3] [169/170] eta: 0:00:00 lr: 0.000050 loss: 0.1918
(0.2092) loss_classifier: 0.0482 (0.0624) loss_box_reg: 0.1245
(0.1286) loss_objectness: 0.0015 (0.0055) loss_rpn_box_reg: 0.0041
(0.0127) time: 0.3639 data: 0.0650 max mem: 2794
Epoch: [3] Total time: 0:01:01 (0.3624 s / it)
mAP:0.452
AP at IoU level [0.50]: 0.721
AP at IoU level [0.55]: 0.706
AP at IoU level [0.60]: 0.678
AP at IoU level [0.65]: 0.653
AP at IoU level [0.70]: 0.641
AP at IoU level [0.75]: 0.513
AP at IoU level [0.80]: 0.424
AP at IoU level [0.85]: 0.162
AP at IoU level [0.90]: 0.018
AP at IoU level [0.95]: 0.000
Epoch: [4] [0/170] eta: 0:00:58 lr: 0.000025 loss: 0.1469
(0.1469) loss_classifier: 0.0314 (0.0314) loss_box_reg: 0.0915
(0.0915) loss_objectness: 0.0083 (0.0083) loss_rpn_box_reg: 0.0157
(0.0157) time: 0.3432 data: 0.0558 max mem: 2794
Epoch: [4] [20/170] eta: 0:00:55 lr: 0.000025 loss: 0.1493
(0.1583) loss_classifier: 0.0425 (0.0469) loss_box_reg: 0.0925
(0.0976) loss_objectness: 0.0022 (0.0038) loss_rpn_box_reg: 0.0037
(0.0100) time: 0.3700 data: 0.0624 max mem: 2794
Epoch: [4] [40/170] eta: 0:00:47 lr: 0.000025 loss: 0.1700
(0.1665) loss_classifier: 0.0458 (0.0494) loss_box_reg: 0.1185
(0.1027) loss_objectness: 0.0011 (0.0038) loss_rpn_box_reg: 0.0038
(0.0106) time: 0.3647 data: 0.0596 max mem: 2794
Epoch: [4] [60/170] eta: 0:00:40 lr: 0.000025 loss: 0.1507
(0.1669) loss_classifier: 0.0484 (0.0489) loss_box_reg: 0.0991
(0.1038) loss_objectness: 0.0007 (0.0036) loss_rpn_box_reg: 0.0028
(0.0107) time: 0.3637 data: 0.0576 max mem: 2794
Epoch: [4] [80/170] eta: 0:00:32 lr: 0.000025 loss: 0.1570
(0.1703) loss_classifier: 0.0456 (0.0493) loss_box_reg: 0.0899
(0.1063) loss_objectness: 0.0023 (0.0037) loss_rpn_box_reg: 0.0057
(0.0110) time: 0.3612 data: 0.0647 max mem: 2794
Epoch: [4] [100/170] eta: 0:00:25 lr: 0.000025 loss: 0.1572
(0.1695) loss_classifier: 0.0417 (0.0492) loss_box_reg: 0.0921
(0.1057) loss_objectness: 0.0008 (0.0042) loss_rpn_box_reg: 0.0033
(0.0104) time: 0.3598 data: 0.0633 max mem: 2794
Epoch: [4] [120/170] eta: 0:00:18 lr: 0.000025 loss: 0.1439
(0.1689) loss_classifier: 0.0473 (0.0499) loss_box_reg: 0.0841
(0.1040) loss_objectness: 0.0006 (0.0039) loss_rpn_box_reg: 0.0029
(0.0112) time: 0.3686 data: 0.0602 max mem: 2794
Epoch: [4] [140/170] eta: 0:00:10 lr: 0.000025 loss: 0.1519
(0.1677) loss_classifier: 0.0419 (0.0490) loss_box_reg: 0.1027
(0.1042) loss_objectness: 0.0014 (0.0037) loss_rpn_box_reg: 0.0058
(0.0109) time: 0.3557 data: 0.0567 max mem: 2794
Epoch: [4] [160/170] eta: 0:00:03 lr: 0.000025 loss: 0.1326

(0.1675) loss_classifier: 0.0358 (0.0485) loss_box_reg: 0.0873
(0.1042) loss_objectness: 0.0013 (0.0039) loss_rpn_box_reg: 0.0050
(0.0109) time: 0.3500 data: 0.0631 max mem: 2794
Epoch: [4] [169/170] eta: 0:00:00 lr: 0.000025 loss: 0.1396
(0.1679) loss_classifier: 0.0391 (0.0486) loss_box_reg: 0.0962
(0.1047) loss_objectness: 0.0012 (0.0039) loss_rpn_box_reg: 0.0034
(0.0108) time: 0.3525 data: 0.0678 max mem: 2794
Epoch: [4] Total time: 0:01:01 (0.3612 s / it)
mAP:0.464
AP at IoU level [0.50]: 0.713
AP at IoU level [0.55]: 0.706
AP at IoU level [0.60]: 0.693
AP at IoU level [0.65]: 0.659
AP at IoU level [0.70]: 0.636
AP at IoU level [0.75]: 0.526
AP at IoU level [0.80]: 0.443
AP at IoU level [0.85]: 0.206
AP at IoU level [0.90]: 0.055
AP at IoU level [0.95]: 0.001
Epoch: [5] [0/170] eta: 0:01:03 lr: 0.000025 loss: 0.1010
(0.1010) loss_classifier: 0.0321 (0.0321) loss_box_reg: 0.0674
(0.0674) loss_objectness: 0.0003 (0.0003) loss_rpn_box_reg: 0.0013
(0.0013) time: 0.3727 data: 0.0593 max mem: 2794
Epoch: [5] [20/170] eta: 0:00:53 lr: 0.000025 loss: 0.1438
(0.1488) loss_classifier: 0.0414 (0.0437) loss_box_reg: 0.0904
(0.0908) loss_objectness: 0.0013 (0.0040) loss_rpn_box_reg: 0.0024
(0.0103) time: 0.3576 data: 0.0631 max mem: 2794
Epoch: [5] [40/170] eta: 0:00:46 lr: 0.000025 loss: 0.1490
(0.1590) loss_classifier: 0.0439 (0.0456) loss_box_reg: 0.1017
(0.0978) loss_objectness: 0.0010 (0.0047) loss_rpn_box_reg: 0.0032
(0.0109) time: 0.3626 data: 0.0618 max mem: 2794
Epoch: [5] [60/170] eta: 0:00:39 lr: 0.000025 loss: 0.1295
(0.1487) loss_classifier: 0.0365 (0.0431) loss_box_reg: 0.0722
(0.0921) loss_objectness: 0.0005 (0.0038) loss_rpn_box_reg: 0.0019
(0.0097) time: 0.3575 data: 0.0562 max mem: 2794
Epoch: [5] [80/170] eta: 0:00:32 lr: 0.000025 loss: 0.1244
(0.1461) loss_classifier: 0.0373 (0.0422) loss_box_reg: 0.0751
(0.0907) loss_objectness: 0.0004 (0.0035) loss_rpn_box_reg: 0.0033
(0.0097) time: 0.3772 data: 0.0607 max mem: 2794
Epoch: [5] [100/170] eta: 0:00:25 lr: 0.000025 loss: 0.1151
(0.1442) loss_classifier: 0.0351 (0.0417) loss_box_reg: 0.0781
(0.0899) loss_objectness: 0.0014 (0.0033) loss_rpn_box_reg: 0.0030
(0.0093) time: 0.3751 data: 0.0667 max mem: 2794
Epoch: [5] [120/170] eta: 0:00:18 lr: 0.000025 loss: 0.1620
(0.1469) loss_classifier: 0.0380 (0.0421) loss_box_reg: 0.1007
(0.0912) loss_objectness: 0.0014 (0.0034) loss_rpn_box_reg: 0.0059
(0.0103) time: 0.3650 data: 0.0598 max mem: 2794
Epoch: [5] [140/170] eta: 0:00:10 lr: 0.000025 loss: 0.1413
(0.1453) loss_classifier: 0.0416 (0.0415) loss_box_reg: 0.0846
(0.0906) loss_objectness: 0.0011 (0.0031) loss_rpn_box_reg: 0.0037

(0.0100) time: 0.3591 data: 0.0564 max mem: 2794
Epoch: [5] [160/170] eta: 0:00:03 lr: 0.000025 loss: 0.1331
(0.1453) loss_classifier: 0.0421 (0.0415) loss_box_reg: 0.0821
(0.0909) loss_objectness: 0.0011 (0.0030) loss_rpn_box_reg: 0.0038
(0.0098) time: 0.3589 data: 0.0617 max mem: 2794
Epoch: [5] [169/170] eta: 0:00:00 lr: 0.000025 loss: 0.1167
(0.1443) loss_classifier: 0.0351 (0.0412) loss_box_reg: 0.0795
(0.0905) loss_objectness: 0.0008 (0.0030) loss_rpn_box_reg: 0.0027
(0.0097) time: 0.3661 data: 0.0622 max mem: 2794
Epoch: [5] Total time: 0:01:01 (0.3645 s / it)
mAP:0.466
AP at IoU level [0.50]: 0.708
AP at IoU level [0.55]: 0.708
AP at IoU level [0.60]: 0.681
AP at IoU level [0.65]: 0.662
AP at IoU level [0.70]: 0.640
AP at IoU level [0.75]: 0.556
AP at IoU level [0.80]: 0.417
AP at IoU level [0.85]: 0.260
AP at IoU level [0.90]: 0.026
AP at IoU level [0.95]: 0.000
Epoch: [6] [0/170] eta: 0:01:02 lr: 0.000013 loss: 0.0843
(0.0843) loss_classifier: 0.0194 (0.0194) loss_box_reg: 0.0589
(0.0589) loss_objectness: 0.0034 (0.0034) loss_rpn_box_reg: 0.0026
(0.0026) time: 0.3689 data: 0.0560 max mem: 2794
Epoch: [6] [20/170] eta: 0:00:53 lr: 0.000013 loss: 0.1094
(0.1265) loss_classifier: 0.0352 (0.0348) loss_box_reg: 0.0653
(0.0788) loss_objectness: 0.0009 (0.0036) loss_rpn_box_reg: 0.0028
(0.0092) time: 0.3589 data: 0.0652 max mem: 2794
Epoch: [6] [40/170] eta: 0:00:46 lr: 0.000013 loss: 0.1089
(0.1248) loss_classifier: 0.0313 (0.0359) loss_box_reg: 0.0672
(0.0756) loss_objectness: 0.0005 (0.0041) loss_rpn_box_reg: 0.0016
(0.0091) time: 0.3590 data: 0.0612 max mem: 2794
Epoch: [6] [60/170] eta: 0:00:39 lr: 0.000013 loss: 0.1017
(0.1232) loss_classifier: 0.0351 (0.0368) loss_box_reg: 0.0616
(0.0751) loss_objectness: 0.0006 (0.0033) loss_rpn_box_reg: 0.0023
(0.0080) time: 0.3597 data: 0.0622 max mem: 2794
Epoch: [6] [80/170] eta: 0:00:32 lr: 0.000013 loss: 0.0913
(0.1189) loss_classifier: 0.0286 (0.0356) loss_box_reg: 0.0616
(0.0735) loss_objectness: 0.0003 (0.0027) loss_rpn_box_reg: 0.0019
(0.0072) time: 0.3677 data: 0.0606 max mem: 2794
Epoch: [6] [100/170] eta: 0:00:25 lr: 0.000013 loss: 0.1161
(0.1207) loss_classifier: 0.0313 (0.0359) loss_box_reg: 0.0726
(0.0747) loss_objectness: 0.0009 (0.0028) loss_rpn_box_reg: 0.0038
(0.0074) time: 0.3681 data: 0.0658 max mem: 2794
Epoch: [6] [120/170] eta: 0:00:18 lr: 0.000013 loss: 0.1258
(0.1220) loss_classifier: 0.0343 (0.0358) loss_box_reg: 0.0771
(0.0758) loss_objectness: 0.0009 (0.0029) loss_rpn_box_reg: 0.0031
(0.0075) time: 0.3541 data: 0.0599 max mem: 2794
Epoch: [6] [140/170] eta: 0:00:10 lr: 0.000013 loss: 0.0863

(0.1204) loss_classifier: 0.0266 (0.0355) loss_box_reg: 0.0585
(0.0748) loss_objectness: 0.0003 (0.0028) loss_rpn_box_reg: 0.0016
(0.0073) time: 0.3559 data: 0.0570 max mem: 2794
Epoch: [6] [160/170] eta: 0:00:03 lr: 0.000013 loss: 0.1283
(0.1226) loss_classifier: 0.0345 (0.0358) loss_box_reg: 0.0864
(0.0758) loss_objectness: 0.0017 (0.0029) loss_rpn_box_reg: 0.0040
(0.0080) time: 0.3624 data: 0.0589 max mem: 2794
Epoch: [6] [169/170] eta: 0:00:00 lr: 0.000013 loss: 0.1326
(0.1234) loss_classifier: 0.0376 (0.0357) loss_box_reg: 0.0864
(0.0766) loss_objectness: 0.0011 (0.0029) loss_rpn_box_reg: 0.0054
(0.0083) time: 0.3681 data: 0.0604 max mem: 2794
Epoch: [6] Total time: 0:01:01 (0.3610 s / it)
mAP:0.472
AP at IoU level [0.50]: 0.707
AP at IoU level [0.55]: 0.699
AP at IoU level [0.60]: 0.679
AP at IoU level [0.65]: 0.661
AP at IoU level [0.70]: 0.645
AP at IoU level [0.75]: 0.569
AP at IoU level [0.80]: 0.438
AP at IoU level [0.85]: 0.266
AP at IoU level [0.90]: 0.054
AP at IoU level [0.95]: 0.001
Epoch: [7] [0/170] eta: 0:01:03 lr: 0.000013 loss: 0.0580
(0.0580) loss_classifier: 0.0123 (0.0123) loss_box_reg: 0.0403
(0.0403) loss_objectness: 0.0014 (0.0014) loss_rpn_box_reg: 0.0040
(0.0040) time: 0.3730 data: 0.0612 max mem: 2794
Epoch: [7] [20/170] eta: 0:00:54 lr: 0.000013 loss: 0.1260
(0.1335) loss_classifier: 0.0392 (0.0399) loss_box_reg: 0.0766
(0.0802) loss_objectness: 0.0008 (0.0037) loss_rpn_box_reg: 0.0036
(0.0096) time: 0.3628 data: 0.0623 max mem: 2794
Epoch: [7] [40/170] eta: 0:00:46 lr: 0.000013 loss: 0.1007
(0.1231) loss_classifier: 0.0310 (0.0364) loss_box_reg: 0.0667
(0.0741) loss_objectness: 0.0003 (0.0037) loss_rpn_box_reg: 0.0015
(0.0089) time: 0.3548 data: 0.0584 max mem: 2794
Epoch: [7] [60/170] eta: 0:00:39 lr: 0.000013 loss: 0.0962
(0.1178) loss_classifier: 0.0272 (0.0352) loss_box_reg: 0.0595
(0.0720) loss_objectness: 0.0008 (0.0029) loss_rpn_box_reg: 0.0033
(0.0077) time: 0.3710 data: 0.0613 max mem: 2794
Epoch: [7] [80/170] eta: 0:00:32 lr: 0.000013 loss: 0.0842
(0.1133) loss_classifier: 0.0277 (0.0341) loss_box_reg: 0.0560
(0.0688) loss_objectness: 0.0003 (0.0026) loss_rpn_box_reg: 0.0024
(0.0078) time: 0.3604 data: 0.0605 max mem: 2794
Epoch: [7] [100/170] eta: 0:00:25 lr: 0.000013 loss: 0.0738
(0.1078) loss_classifier: 0.0224 (0.0325) loss_box_reg: 0.0495
(0.0654) loss_objectness: 0.0002 (0.0025) loss_rpn_box_reg: 0.0012
(0.0075) time: 0.3573 data: 0.0630 max mem: 2794
Epoch: [7] [120/170] eta: 0:00:18 lr: 0.000013 loss: 0.0805
(0.1078) loss_classifier: 0.0290 (0.0321) loss_box_reg: 0.0508
(0.0650) loss_objectness: 0.0006 (0.0026) loss_rpn_box_reg: 0.0023

(0.0081) time: 0.3596 data: 0.0590 max mem: 2794
Epoch: [7] [140/170] eta: 0:00:10 lr: 0.000013 loss: 0.0982
(0.1079) loss_classifier: 0.0294 (0.0324) loss_box_reg: 0.0628
(0.0651) loss_objectness: 0.0007 (0.0025) loss_rpn_box_reg: 0.0028
(0.0079) time: 0.3655 data: 0.0609 max mem: 2794
Epoch: [7] [160/170] eta: 0:00:03 lr: 0.000013 loss: 0.0855
(0.1084) loss_classifier: 0.0284 (0.0326) loss_box_reg: 0.0480
(0.0654) loss_objectness: 0.0006 (0.0026) loss_rpn_box_reg: 0.0016
(0.0079) time: 0.3530 data: 0.0588 max mem: 2794
Epoch: [7] [169/170] eta: 0:00:00 lr: 0.000013 loss: 0.0926
(0.1076) loss_classifier: 0.0289 (0.0324) loss_box_reg: 0.0594
(0.0649) loss_objectness: 0.0014 (0.0026) loss_rpn_box_reg: 0.0039
(0.0077) time: 0.3675 data: 0.0628 max mem: 2794
Epoch: [7] Total time: 0:01:01 (0.3620 s / it)
mAP:0.464
AP at IoU level [0.50]: 0.717
AP at IoU level [0.55]: 0.710
AP at IoU level [0.60]: 0.688
AP at IoU level [0.65]: 0.671
AP at IoU level [0.70]: 0.632
AP at IoU level [0.75]: 0.544
AP at IoU level [0.80]: 0.376
AP at IoU level [0.85]: 0.258
AP at IoU level [0.90]: 0.038
AP at IoU level [0.95]: 0.000
Epoch: [8] [0/170] eta: 0:00:57 lr: 0.000006 loss: 0.0945
(0.0945) loss_classifier: 0.0234 (0.0234) loss_box_reg: 0.0614
(0.0614) loss_objectness: 0.0010 (0.0010) loss_rpn_box_reg: 0.0087
(0.0087) time: 0.3396 data: 0.0516 max mem: 2794
Epoch: [8] [20/170] eta: 0:00:54 lr: 0.000006 loss: 0.0794
(0.1004) loss_classifier: 0.0319 (0.0300) loss_box_reg: 0.0472
(0.0567) loss_objectness: 0.0006 (0.0044) loss_rpn_box_reg: 0.0035
(0.0094) time: 0.3649 data: 0.0588 max mem: 2794
Epoch: [8] [40/170] eta: 0:00:46 lr: 0.000006 loss: 0.0817
(0.0969) loss_classifier: 0.0274 (0.0300) loss_box_reg: 0.0477
(0.0558) loss_objectness: 0.0002 (0.0039) loss_rpn_box_reg: 0.0014
(0.0072) time: 0.3453 data: 0.0618 max mem: 2794
Epoch: [8] [60/170] eta: 0:00:39 lr: 0.000006 loss: 0.0966
(0.1053) loss_classifier: 0.0304 (0.0314) loss_box_reg: 0.0616
(0.0616) loss_objectness: 0.0003 (0.0034) loss_rpn_box_reg: 0.0033
(0.0088) time: 0.3681 data: 0.0591 max mem: 2794
Epoch: [8] [80/170] eta: 0:00:32 lr: 0.000006 loss: 0.0829
(0.1030) loss_classifier: 0.0280 (0.0309) loss_box_reg: 0.0500
(0.0609) loss_objectness: 0.0003 (0.0032) loss_rpn_box_reg: 0.0023
(0.0080) time: 0.3675 data: 0.0640 max mem: 2794
Epoch: [8] [100/170] eta: 0:00:25 lr: 0.000006 loss: 0.0624
(0.0971) loss_classifier: 0.0212 (0.0295) loss_box_reg: 0.0381
(0.0576) loss_objectness: 0.0006 (0.0028) loss_rpn_box_reg: 0.0017
(0.0072) time: 0.3708 data: 0.0624 max mem: 2794
Epoch: [8] [120/170] eta: 0:00:18 lr: 0.000006 loss: 0.0732

```
(0.0947) loss_classifier: 0.0268 (0.0293) loss_box_reg: 0.0396
(0.0559) loss_objectness: 0.0003 (0.0025) loss_rpn_box_reg: 0.0023
(0.0070) time: 0.3544 data: 0.0595 max mem: 2794
Epoch: [8] [140/170] eta: 0:00:10 lr: 0.000006 loss: 0.0844
(0.0950) loss_classifier: 0.0267 (0.0295) loss_box_reg: 0.0484
(0.0559) loss_objectness: 0.0004 (0.0025) loss_rpn_box_reg: 0.0020
(0.0071) time: 0.3534 data: 0.0599 max mem: 2794
Epoch: [8] [160/170] eta: 0:00:03 lr: 0.000006 loss: 0.0880
(0.0947) loss_classifier: 0.0274 (0.0295) loss_box_reg: 0.0532
(0.0561) loss_objectness: 0.0002 (0.0023) loss_rpn_box_reg: 0.0028
(0.0068) time: 0.3597 data: 0.0577 max mem: 2794
Epoch: [8] [169/170] eta: 0:00:00 lr: 0.000006 loss: 0.0928
(0.0960) loss_classifier: 0.0276 (0.0296) loss_box_reg: 0.0547
(0.0569) loss_objectness: 0.0010 (0.0024) loss_rpn_box_reg: 0.0035
(0.0071) time: 0.3558 data: 0.0588 max mem: 2794
Epoch: [8] Total time: 0:01:01 (0.3605 s / it)
mAP:0.462
  AP at IoU level [0.50]: 0.715
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.678
  AP at IoU level [0.65]: 0.662
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.549
  AP at IoU level [0.80]: 0.378
  AP at IoU level [0.85]: 0.261
  AP at IoU level [0.90]: 0.049
  AP at IoU level [0.95]: 0.001
Epoch: [9] [ 0/170] eta: 0:01:02 lr: 0.000006 loss: 0.1462
(0.1462) loss_classifier: 0.0343 (0.0343) loss_box_reg: 0.1022
(0.1022) loss_objectness: 0.0014 (0.0014) loss_rpn_box_reg: 0.0083
(0.0083) time: 0.3681 data: 0.0603 max mem: 2794
Epoch: [9] [ 20/170] eta: 0:00:53 lr: 0.000006 loss: 0.0819
(0.0990) loss_classifier: 0.0294 (0.0332) loss_box_reg: 0.0502
(0.0579) loss_objectness: 0.0003 (0.0020) loss_rpn_box_reg: 0.0016
(0.0059) time: 0.3577 data: 0.0604 max mem: 2794
Epoch: [9] [ 40/170] eta: 0:00:46 lr: 0.000006 loss: 0.0554
(0.0850) loss_classifier: 0.0205 (0.0287) loss_box_reg: 0.0311
(0.0494) loss_objectness: 0.0001 (0.0015) loss_rpn_box_reg: 0.0006
(0.0054) time: 0.3616 data: 0.0577 max mem: 2794
Epoch: [9] [ 60/170] eta: 0:00:39 lr: 0.000006 loss: 0.0893
(0.0863) loss_classifier: 0.0258 (0.0290) loss_box_reg: 0.0508
(0.0497) loss_objectness: 0.0006 (0.0017) loss_rpn_box_reg: 0.0056
(0.0059) time: 0.3516 data: 0.0633 max mem: 2794
Epoch: [9] [ 80/170] eta: 0:00:32 lr: 0.000006 loss: 0.0631
(0.0871) loss_classifier: 0.0207 (0.0284) loss_box_reg: 0.0408
(0.0505) loss_objectness: 0.0006 (0.0020) loss_rpn_box_reg: 0.0023
(0.0063) time: 0.3567 data: 0.0586 max mem: 2794
Epoch: [9] [100/170] eta: 0:00:25 lr: 0.000006 loss: 0.0671
(0.0891) loss_classifier: 0.0250 (0.0287) loss_box_reg: 0.0398
(0.0511) loss_objectness: 0.0005 (0.0023) loss_rpn_box_reg: 0.0012
```

```
(0.0070) time: 0.3680 data: 0.0610 max mem: 2794
Epoch: [9] [120/170] eta: 0:00:18 lr: 0.000006 loss: 0.0619
(0.0858) loss_classifier: 0.0222 (0.0276) loss_box_reg: 0.0387
(0.0498) loss_objectness: 0.0002 (0.0021) loss_rpn_box_reg: 0.0011
(0.0063) time: 0.3646 data: 0.0604 max mem: 2794
Epoch: [9] [140/170] eta: 0:00:10 lr: 0.000006 loss: 0.0866
(0.0877) loss_classifier: 0.0272 (0.0284) loss_box_reg: 0.0535
(0.0509) loss_objectness: 0.0008 (0.0021) loss_rpn_box_reg: 0.0030
(0.0064) time: 0.3612 data: 0.0615 max mem: 2794
Epoch: [9] [160/170] eta: 0:00:03 lr: 0.000006 loss: 0.0665
(0.0887) loss_classifier: 0.0296 (0.0286) loss_box_reg: 0.0404
(0.0512) loss_objectness: 0.0001 (0.0023) loss_rpn_box_reg: 0.0018
(0.0067) time: 0.3473 data: 0.0573 max mem: 2794
Epoch: [9] [169/170] eta: 0:00:00 lr: 0.000006 loss: 0.0628
(0.0898) loss_classifier: 0.0275 (0.0288) loss_box_reg: 0.0335
(0.0518) loss_objectness: 0.0001 (0.0024) loss_rpn_box_reg: 0.0013
(0.0067) time: 0.3523 data: 0.0573 max mem: 2794
Epoch: [9] Total time: 0:01:00 (0.3585 s / it)
mAP:0.474
  AP at IoU level [0.50]: 0.715
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.677
  AP at IoU level [0.65]: 0.669
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.578
  AP at IoU level [0.80]: 0.452
  AP at IoU level [0.85]: 0.261
  AP at IoU level [0.90]: 0.061
  AP at IoU level [0.95]: 0.002
Epoch: [10] [ 0/170] eta: 0:01:03 lr: 0.000003 loss: 0.0833
(0.0833) loss_classifier: 0.0289 (0.0289) loss_box_reg: 0.0525
(0.0525) loss_objectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0018
(0.0018) time: 0.3711 data: 0.0585 max mem: 2794
Epoch: [10] [ 20/170] eta: 0:00:53 lr: 0.000003 loss: 0.0820
(0.0726) loss_classifier: 0.0228 (0.0241) loss_box_reg: 0.0478
(0.0437) loss_objectness: 0.0005 (0.0011) loss_rpn_box_reg: 0.0031
(0.0037) time: 0.3570 data: 0.0581 max mem: 2794
Epoch: [10] [ 40/170] eta: 0:00:46 lr: 0.000003 loss: 0.0582
(0.0752) loss_classifier: 0.0243 (0.0254) loss_box_reg: 0.0303
(0.0432) loss_objectness: 0.0001 (0.0018) loss_rpn_box_reg: 0.0011
(0.0049) time: 0.3636 data: 0.0614 max mem: 2794
Epoch: [10] [ 60/170] eta: 0:00:39 lr: 0.000003 loss: 0.0898
(0.0811) loss_classifier: 0.0289 (0.0268) loss_box_reg: 0.0524
(0.0462) loss_objectness: 0.0007 (0.0020) loss_rpn_box_reg: 0.0051
(0.0061) time: 0.3648 data: 0.0603 max mem: 2794
Epoch: [10] [ 80/170] eta: 0:00:32 lr: 0.000003 loss: 0.0747
(0.0821) loss_classifier: 0.0263 (0.0277) loss_box_reg: 0.0466
(0.0467) loss_objectness: 0.0003 (0.0018) loss_rpn_box_reg: 0.0022
(0.0059) time: 0.3589 data: 0.0623 max mem: 2794
Epoch: [10] [100/170] eta: 0:00:25 lr: 0.000003 loss: 0.0841
```

```
(0.0845) loss_classifier: 0.0293 (0.0284) loss_box_reg: 0.0511
(0.0477) loss_objectness: 0.0008 (0.0019) loss_rpn_box_reg: 0.0025
(0.0065) time: 0.3758 data: 0.0691 max mem: 2794
Epoch: [10] [120/170] eta: 0:00:18 lr: 0.000003 loss: 0.0620
(0.0863) loss_classifier: 0.0270 (0.0286) loss_box_reg: 0.0354
(0.0488) loss_objectness: 0.0002 (0.0022) loss_rpn_box_reg: 0.0022
(0.0066) time: 0.3605 data: 0.0585 max mem: 2794
Epoch: [10] [140/170] eta: 0:00:10 lr: 0.000003 loss: 0.0519
(0.0836) loss_classifier: 0.0216 (0.0278) loss_box_reg: 0.0263
(0.0475) loss_objectness: 0.0002 (0.0021) loss_rpn_box_reg: 0.0010
(0.0063) time: 0.3464 data: 0.0578 max mem: 2794
Epoch: [10] [160/170] eta: 0:00:03 lr: 0.000003 loss: 0.0636
(0.0848) loss_classifier: 0.0250 (0.0280) loss_box_reg: 0.0392
(0.0480) loss_objectness: 0.0005 (0.0022) loss_rpn_box_reg: 0.0019
(0.0067) time: 0.3644 data: 0.0638 max mem: 2794
Epoch: [10] [169/170] eta: 0:00:00 lr: 0.000003 loss: 0.0594
(0.0837) loss_classifier: 0.0236 (0.0278) loss_box_reg: 0.0333
(0.0473) loss_objectness: 0.0003 (0.0021) loss_rpn_box_reg: 0.0013
(0.0065) time: 0.3542 data: 0.0605 max mem: 2794
Epoch: [10] Total time: 0:01:01 (0.3612 s / it)
mAP:0.459
  AP at IoU level [0.50]: 0.707
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.676
  AP at IoU level [0.65]: 0.651
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.543
  AP at IoU level [0.80]: 0.399
  AP at IoU level [0.85]: 0.242
  AP at IoU level [0.90]: 0.047
  AP at IoU level [0.95]: 0.000
Epoch: [11] [ 0/170] eta: 0:01:10 lr: 0.000003 loss: 0.0280
(0.0280) loss_classifier: 0.0128 (0.0128) loss_box_reg: 0.0148
(0.0148) loss_objectness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0004
(0.0004) time: 0.4151 data: 0.0582 max mem: 2794
Epoch: [11] [ 20/170] eta: 0:00:56 lr: 0.000003 loss: 0.0641
(0.0688) loss_classifier: 0.0221 (0.0244) loss_box_reg: 0.0339
(0.0395) loss_objectness: 0.0002 (0.0014) loss_rpn_box_reg: 0.0020
(0.0036) time: 0.3736 data: 0.0632 max mem: 2795
Epoch: [11] [ 40/170] eta: 0:00:47 lr: 0.000003 loss: 0.0525
(0.0681) loss_classifier: 0.0218 (0.0244) loss_box_reg: 0.0318
(0.0389) loss_objectness: 0.0002 (0.0016) loss_rpn_box_reg: 0.0012
(0.0031) time: 0.3508 data: 0.0614 max mem: 2795
Epoch: [11] [ 60/170] eta: 0:00:39 lr: 0.000003 loss: 0.0757
(0.0736) loss_classifier: 0.0231 (0.0251) loss_box_reg: 0.0435
(0.0418) loss_objectness: 0.0007 (0.0018) loss_rpn_box_reg: 0.0021
(0.0049) time: 0.3618 data: 0.0565 max mem: 2795
Epoch: [11] [ 80/170] eta: 0:00:32 lr: 0.000003 loss: 0.0645
(0.0728) loss_classifier: 0.0253 (0.0248) loss_box_reg: 0.0370
(0.0412) loss_objectness: 0.0003 (0.0018) loss_rpn_box_reg: 0.0017
```

```
(0.0049) time: 0.3633 data: 0.0597 max mem: 2795
Epoch: [11] [100/170] eta: 0:00:25 lr: 0.000003 loss: 0.0882
(0.0788) loss_classifier: 0.0285 (0.0263) loss_box_reg: 0.0491
(0.0444) loss_objectness: 0.0009 (0.0020) loss_rpn_box_reg: 0.0044
(0.0060) time: 0.3744 data: 0.0665 max mem: 2795
Epoch: [11] [120/170] eta: 0:00:18 lr: 0.000003 loss: 0.0608
(0.0776) loss_classifier: 0.0210 (0.0261) loss_box_reg: 0.0338
(0.0439) loss_objectness: 0.0002 (0.0018) loss_rpn_box_reg: 0.0012
(0.0058) time: 0.3498 data: 0.0596 max mem: 2795
Epoch: [11] [140/170] eta: 0:00:10 lr: 0.000003 loss: 0.0481
(0.0769) loss_classifier: 0.0210 (0.0259) loss_box_reg: 0.0240
(0.0433) loss_objectness: 0.0001 (0.0018) loss_rpn_box_reg: 0.0010
(0.0059) time: 0.3678 data: 0.0596 max mem: 2795
Epoch: [11] [160/170] eta: 0:00:03 lr: 0.000003 loss: 0.0725
(0.0788) loss_classifier: 0.0275 (0.0265) loss_box_reg: 0.0436
(0.0443) loss_objectness: 0.0004 (0.0020) loss_rpn_box_reg: 0.0024
(0.0060) time: 0.3500 data: 0.0623 max mem: 2795
Epoch: [11] [169/170] eta: 0:00:00 lr: 0.000003 loss: 0.0823
(0.0801) loss_classifier: 0.0295 (0.0267) loss_box_reg: 0.0474
(0.0449) loss_objectness: 0.0007 (0.0021) loss_rpn_box_reg: 0.0030
(0.0063) time: 0.3546 data: 0.0614 max mem: 2795
Epoch: [11] Total time: 0:01:01 (0.3617 s / it)
mAP:0.467
  AP at IoU level [0.50]: 0.715
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.676
  AP at IoU level [0.65]: 0.668
  AP at IoU level [0.70]: 0.629
  AP at IoU level [0.75]: 0.581
  AP at IoU level [0.80]: 0.420
  AP at IoU level [0.85]: 0.248
  AP at IoU level [0.90]: 0.039
  AP at IoU level [0.95]: 0.000
Epoch: [12] [ 0/170] eta: 0:01:05 lr: 0.000002 loss: 0.0345
(0.0345) loss_classifier: 0.0172 (0.0172) loss_box_reg: 0.0170
(0.0170) loss_objectness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0003
(0.0003) time: 0.3872 data: 0.0527 max mem: 2795
Epoch: [12] [ 20/170] eta: 0:00:55 lr: 0.000002 loss: 0.0652
(0.0682) loss_classifier: 0.0233 (0.0242) loss_box_reg: 0.0377
(0.0367) loss_objectness: 0.0003 (0.0013) loss_rpn_box_reg: 0.0019
(0.0060) time: 0.3686 data: 0.0652 max mem: 2795
Epoch: [12] [ 40/170] eta: 0:00:47 lr: 0.000002 loss: 0.0747
(0.0781) loss_classifier: 0.0271 (0.0264) loss_box_reg: 0.0470
(0.0435) loss_objectness: 0.0008 (0.0014) loss_rpn_box_reg: 0.0038
(0.0067) time: 0.3540 data: 0.0617 max mem: 2795
Epoch: [12] [ 60/170] eta: 0:00:39 lr: 0.000002 loss: 0.0567
(0.0780) loss_classifier: 0.0225 (0.0261) loss_box_reg: 0.0322
(0.0435) loss_objectness: 0.0005 (0.0019) loss_rpn_box_reg: 0.0035
(0.0065) time: 0.3653 data: 0.0565 max mem: 2795
Epoch: [12] [ 80/170] eta: 0:00:32 lr: 0.000002 loss: 0.0824
```



```
(0.0807) loss_classifier: 0.0238 (0.0261) loss_box_reg: 0.0470
(0.0448) loss_objectness: 0.0011 (0.0026) loss_rpn_box_reg: 0.0025
(0.0072) time: 0.3554 data: 0.0614 max mem: 2795
Epoch: [12] [100/170] eta: 0:00:25 lr: 0.000002 loss: 0.0684
(0.0804) loss_classifier: 0.0265 (0.0264) loss_box_reg: 0.0399
(0.0445) loss_objectness: 0.0006 (0.0026) loss_rpn_box_reg: 0.0023
(0.0070) time: 0.3663 data: 0.0683 max mem: 2795
Epoch: [12] [120/170] eta: 0:00:17 lr: 0.000002 loss: 0.0510
(0.0793) loss_classifier: 0.0212 (0.0261) loss_box_reg: 0.0272
(0.0439) loss_objectness: 0.0001 (0.0025) loss_rpn_box_reg: 0.0008
(0.0069) time: 0.3489 data: 0.0570 max mem: 2795
Epoch: [12] [140/170] eta: 0:00:10 lr: 0.000002 loss: 0.0612
(0.0781) loss_classifier: 0.0248 (0.0261) loss_box_reg: 0.0319
(0.0432) loss_objectness: 0.0004 (0.0023) loss_rpn_box_reg: 0.0020
(0.0065) time: 0.3657 data: 0.0614 max mem: 2795
Epoch: [12] [160/170] eta: 0:00:03 lr: 0.000002 loss: 0.0562
(0.0766) loss_classifier: 0.0231 (0.0259) loss_box_reg: 0.0311
(0.0423) loss_objectness: 0.0001 (0.0022) loss_rpn_box_reg: 0.0011
(0.0062) time: 0.3522 data: 0.0588 max mem: 2795
Epoch: [12] [169/170] eta: 0:00:00 lr: 0.000002 loss: 0.0582
(0.0766) loss_classifier: 0.0201 (0.0258) loss_box_reg: 0.0325
(0.0423) loss_objectness: 0.0003 (0.0023) loss_rpn_box_reg: 0.0014
(0.0062) time: 0.3541 data: 0.0634 max mem: 2795
Epoch: [12] Total time: 0:01:01 (0.3598 s / it)
mAP:0.466
  AP at IoU level [0.50]: 0.715
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.677
  AP at IoU level [0.65]: 0.669
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.574
  AP at IoU level [0.80]: 0.402
  AP at IoU level [0.85]: 0.246
  AP at IoU level [0.90]: 0.048
  AP at IoU level [0.95]: 0.000
Epoch: [13] [ 0/170] eta: 0:01:05 lr: 0.000002 loss: 0.0385
(0.0385) loss_classifier: 0.0164 (0.0164) loss_box_reg: 0.0218
(0.0218) loss_objectness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0002
(0.0002) time: 0.3866 data: 0.0742 max mem: 2795
Epoch: [13] [ 20/170] eta: 0:00:55 lr: 0.000002 loss: 0.0547
(0.0756) loss_classifier: 0.0209 (0.0247) loss_box_reg: 0.0321
(0.0409) loss_objectness: 0.0003 (0.0021) loss_rpn_box_reg: 0.0021
(0.0079) time: 0.3715 data: 0.0614 max mem: 2795
Epoch: [13] [ 40/170] eta: 0:00:48 lr: 0.000002 loss: 0.0516
(0.0726) loss_classifier: 0.0184 (0.0238) loss_box_reg: 0.0282
(0.0402) loss_objectness: 0.0004 (0.0021) loss_rpn_box_reg: 0.0017
(0.0064) time: 0.3687 data: 0.0585 max mem: 2795
Epoch: [13] [ 60/170] eta: 0:00:40 lr: 0.000002 loss: 0.0641
(0.0739) loss_classifier: 0.0250 (0.0246) loss_box_reg: 0.0319
(0.0404) loss_objectness: 0.0007 (0.0020) loss_rpn_box_reg: 0.0017
```

```
(0.0068) time: 0.3598 data: 0.0625 max mem: 2795
Epoch: [13] [ 80/170] eta: 0:00:32 lr: 0.000002 loss: 0.0693
(0.0733) loss_classifier: 0.0265 (0.0248) loss_box_reg: 0.0366
(0.0402) loss_objectness: 0.0001 (0.0018) loss_rpn_box_reg: 0.0017
(0.0065) time: 0.3612 data: 0.0590 max mem: 2795
Epoch: [13] [100/170] eta: 0:00:25 lr: 0.000002 loss: 0.0726
(0.0780) loss_classifier: 0.0259 (0.0264) loss_box_reg: 0.0405
(0.0428) loss_objectness: 0.0008 (0.0019) loss_rpn_box_reg: 0.0051
(0.0069) time: 0.3682 data: 0.0643 max mem: 2795
Epoch: [13] [120/170] eta: 0:00:18 lr: 0.000002 loss: 0.0619
(0.0763) loss_classifier: 0.0202 (0.0258) loss_box_reg: 0.0364
(0.0423) loss_objectness: 0.0003 (0.0019) loss_rpn_box_reg: 0.0015
(0.0064) time: 0.3452 data: 0.0639 max mem: 2795
Epoch: [13] [140/170] eta: 0:00:10 lr: 0.000002 loss: 0.0593
(0.0758) loss_classifier: 0.0218 (0.0256) loss_box_reg: 0.0324
(0.0420) loss_objectness: 0.0004 (0.0019) loss_rpn_box_reg: 0.0026
(0.0063) time: 0.3539 data: 0.0600 max mem: 2795
Epoch: [13] [160/170] eta: 0:00:03 lr: 0.000002 loss: 0.0761
(0.0761) loss_classifier: 0.0281 (0.0259) loss_box_reg: 0.0395
(0.0421) loss_objectness: 0.0003 (0.0018) loss_rpn_box_reg: 0.0024
(0.0063) time: 0.3459 data: 0.0607 max mem: 2795
Epoch: [13] [169/170] eta: 0:00:00 lr: 0.000002 loss: 0.0628
(0.0751) loss_classifier: 0.0251 (0.0257) loss_box_reg: 0.0388
(0.0416) loss_objectness: 0.0005 (0.0018) loss_rpn_box_reg: 0.0025
(0.0061) time: 0.3487 data: 0.0625 max mem: 2795
Epoch: [13] Total time: 0:01:01 (0.3594 s / it)
mAP:0.466
  AP at IoU level [0.50]: 0.715
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.677
  AP at IoU level [0.65]: 0.669
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.575
  AP at IoU level [0.80]: 0.403
  AP at IoU level [0.85]: 0.253
  AP at IoU level [0.90]: 0.036
  AP at IoU level [0.95]: 0.000
Epoch: [14] [ 0/170] eta: 0:01:05 lr: 0.000001 loss: 0.0389
(0.0389) loss_classifier: 0.0208 (0.0208) loss_box_reg: 0.0177
(0.0177) loss_objectness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0003
(0.0003) time: 0.3834 data: 0.0489 max mem: 2795
Epoch: [14] [ 20/170] eta: 0:00:54 lr: 0.000001 loss: 0.0697
(0.0788) loss_classifier: 0.0206 (0.0262) loss_box_reg: 0.0374
(0.0409) loss_objectness: 0.0003 (0.0029) loss_rpn_box_reg: 0.0014
(0.0088) time: 0.3590 data: 0.0658 max mem: 2795
Epoch: [14] [ 40/170] eta: 0:00:47 lr: 0.000001 loss: 0.0545
(0.0739) loss_classifier: 0.0227 (0.0253) loss_box_reg: 0.0285
(0.0400) loss_objectness: 0.0003 (0.0022) loss_rpn_box_reg: 0.0010
(0.0064) time: 0.3638 data: 0.0617 max mem: 2795
Epoch: [14] [ 60/170] eta: 0:00:39 lr: 0.000001 loss: 0.0524
```

(0.0737) loss_classifier: 0.0220 (0.0250) loss_box_reg: 0.0349
(0.0397) loss_objectness: 0.0005 (0.0022) loss_rpn_box_reg: 0.0022
(0.0068) time: 0.3637 data: 0.0567 max mem: 2795
Epoch: [14] [80/170] eta: 0:00:32 lr: 0.000001 loss: 0.0663
(0.0756) loss_classifier: 0.0257 (0.0256) loss_box_reg: 0.0353
(0.0409) loss_objectness: 0.0004 (0.0025) loss_rpn_box_reg: 0.0022
(0.0067) time: 0.3593 data: 0.0626 max mem: 2795
Epoch: [14] [100/170] eta: 0:00:25 lr: 0.000001 loss: 0.0641
(0.0763) loss_classifier: 0.0270 (0.0258) loss_box_reg: 0.0317
(0.0414) loss_objectness: 0.0002 (0.0025) loss_rpn_box_reg: 0.0018
(0.0066) time: 0.3714 data: 0.0657 max mem: 2795
Epoch: [14] [120/170] eta: 0:00:18 lr: 0.000001 loss: 0.0582
(0.0764) loss_classifier: 0.0243 (0.0256) loss_box_reg: 0.0298
(0.0416) loss_objectness: 0.0007 (0.0024) loss_rpn_box_reg: 0.0022
(0.0067) time: 0.3630 data: 0.0616 max mem: 2795
Epoch: [14] [140/170] eta: 0:00:10 lr: 0.000001 loss: 0.0576
(0.0745) loss_classifier: 0.0214 (0.0252) loss_box_reg: 0.0343
(0.0406) loss_objectness: 0.0002 (0.0023) loss_rpn_box_reg: 0.0021
(0.0063) time: 0.3679 data: 0.0593 max mem: 2795
Epoch: [14] [160/170] eta: 0:00:03 lr: 0.000001 loss: 0.0679
(0.0744) loss_classifier: 0.0236 (0.0254) loss_box_reg: 0.0347
(0.0408) loss_objectness: 0.0001 (0.0022) loss_rpn_box_reg: 0.0012
(0.0060) time: 0.3686 data: 0.0625 max mem: 2795
Epoch: [14] [169/170] eta: 0:00:00 lr: 0.000001 loss: 0.0546
(0.0737) loss_classifier: 0.0243 (0.0253) loss_box_reg: 0.0318
(0.0404) loss_objectness: 0.0002 (0.0021) loss_rpn_box_reg: 0.0012
(0.0059) time: 0.3631 data: 0.0601 max mem: 2795
Epoch: [14] Total time: 0:01:01 (0.3642 s / it)
mAP:0.469

AP at IoU level [0.50]: 0.715
AP at IoU level [0.55]: 0.698
AP at IoU level [0.60]: 0.677
AP at IoU level [0.65]: 0.669
AP at IoU level [0.70]: 0.632
AP at IoU level [0.75]: 0.575
AP at IoU level [0.80]: 0.423
AP at IoU level [0.85]: 0.260
AP at IoU level [0.90]: 0.041
AP at IoU level [0.95]: 0.000

Epoch: [15] [0/170] eta: 0:01:01 lr: 0.000001 loss: 0.0404
(0.0404) loss_classifier: 0.0167 (0.0167) loss_box_reg: 0.0226
(0.0226) loss_objectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0010
(0.0010) time: 0.3606 data: 0.0520 max mem: 2795
Epoch: [15] [20/170] eta: 0:00:53 lr: 0.000001 loss: 0.0383
(0.0671) loss_classifier: 0.0174 (0.0226) loss_box_reg: 0.0220
(0.0373) loss_objectness: 0.0001 (0.0026) loss_rpn_box_reg: 0.0008
(0.0047) time: 0.3559 data: 0.0603 max mem: 2795
Epoch: [15] [40/170] eta: 0:00:46 lr: 0.000001 loss: 0.0706
(0.0782) loss_classifier: 0.0221 (0.0255) loss_box_reg: 0.0349
(0.0429) loss_objectness: 0.0007 (0.0028) loss_rpn_box_reg: 0.0014

```
(0.0070) time: 0.3663 data: 0.0646 max mem: 2795
Epoch: [15] [ 60/170] eta: 0:00:39 lr: 0.000001 loss: 0.0605
(0.0749) loss_classifier: 0.0254 (0.0256) loss_box_reg: 0.0292
(0.0411) loss_objectness: 0.0003 (0.0022) loss_rpn_box_reg: 0.0014
(0.0060) time: 0.3611 data: 0.0598 max mem: 2795
Epoch: [15] [ 80/170] eta: 0:00:32 lr: 0.000001 loss: 0.0648
(0.0739) loss_classifier: 0.0224 (0.0254) loss_box_reg: 0.0343
(0.0404) loss_objectness: 0.0003 (0.0021) loss_rpn_box_reg: 0.0023
(0.0060) time: 0.3617 data: 0.0601 max mem: 2795
Epoch: [15] [100/170] eta: 0:00:25 lr: 0.000001 loss: 0.0710
(0.0733) loss_classifier: 0.0234 (0.0252) loss_box_reg: 0.0382
(0.0403) loss_objectness: 0.0005 (0.0020) loss_rpn_box_reg: 0.0048
(0.0058) time: 0.3719 data: 0.0577 max mem: 2795
Epoch: [15] [120/170] eta: 0:00:18 lr: 0.000001 loss: 0.0666
(0.0744) loss_classifier: 0.0256 (0.0256) loss_box_reg: 0.0358
(0.0409) loss_objectness: 0.0005 (0.0020) loss_rpn_box_reg: 0.0016
(0.0059) time: 0.3721 data: 0.0706 max mem: 2795
Epoch: [15] [140/170] eta: 0:00:10 lr: 0.000001 loss: 0.0715
(0.0768) loss_classifier: 0.0240 (0.0260) loss_box_reg: 0.0371
(0.0419) loss_objectness: 0.0008 (0.0024) loss_rpn_box_reg: 0.0030
(0.0066) time: 0.3682 data: 0.0632 max mem: 2795
Epoch: [15] [160/170] eta: 0:00:03 lr: 0.000001 loss: 0.0545
(0.0747) loss_classifier: 0.0203 (0.0257) loss_box_reg: 0.0287
(0.0407) loss_objectness: 0.0002 (0.0022) loss_rpn_box_reg: 0.0011
(0.0062) time: 0.3641 data: 0.0647 max mem: 2795
Epoch: [15] [169/170] eta: 0:00:00 lr: 0.000001 loss: 0.0639
(0.0738) loss_classifier: 0.0227 (0.0254) loss_box_reg: 0.0327
(0.0401) loss_objectness: 0.0002 (0.0021) loss_rpn_box_reg: 0.0020
(0.0061) time: 0.3647 data: 0.0611 max mem: 2795
Epoch: [15] Total time: 0:01:02 (0.3651 s / it)
mAP:0.465
  AP at IoU level [0.50]: 0.714
  AP at IoU level [0.55]: 0.697
  AP at IoU level [0.60]: 0.676
  AP at IoU level [0.65]: 0.668
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.574
  AP at IoU level [0.80]: 0.402
  AP at IoU level [0.85]: 0.243
  AP at IoU level [0.90]: 0.046
  AP at IoU level [0.95]: 0.000
Epoch: [16] [ 0/170] eta: 0:01:05 lr: 0.000000 loss: 0.0529
(0.0529) loss_classifier: 0.0243 (0.0243) loss_box_reg: 0.0276
(0.0276) loss_objectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0010
(0.0010) time: 0.3842 data: 0.0983 max mem: 2795
Epoch: [16] [ 20/170] eta: 0:00:55 lr: 0.000000 loss: 0.0448
(0.0585) loss_classifier: 0.0173 (0.0217) loss_box_reg: 0.0258
(0.0317) loss_objectness: 0.0002 (0.0009) loss_rpn_box_reg: 0.0010
(0.0042) time: 0.3703 data: 0.0588 max mem: 2795
Epoch: [16] [ 40/170] eta: 0:00:47 lr: 0.000000 loss: 0.0896
```

```
(0.0801) loss_classifier: 0.0294 (0.0267) loss_box_reg: 0.0441
(0.0430) loss_objectness: 0.0011 (0.0029) loss_rpn_box_reg: 0.0038
(0.0075) time: 0.3576 data: 0.0603 max mem: 2795
Epoch: [16] [ 60/170] eta: 0:00:39 lr: 0.000000 loss: 0.0672
(0.0772) loss_classifier: 0.0244 (0.0260) loss_box_reg: 0.0336
(0.0417) loss_objectness: 0.0010 (0.0025) loss_rpn_box_reg: 0.0029
(0.0070) time: 0.3603 data: 0.0586 max mem: 2795
Epoch: [16] [ 80/170] eta: 0:00:32 lr: 0.000000 loss: 0.0510
(0.0739) loss_classifier: 0.0177 (0.0254) loss_box_reg: 0.0293
(0.0398) loss_objectness: 0.0002 (0.0021) loss_rpn_box_reg: 0.0013
(0.0066) time: 0.3655 data: 0.0619 max mem: 2795
Epoch: [16] [100/170] eta: 0:00:25 lr: 0.000000 loss: 0.0602
(0.0745) loss_classifier: 0.0214 (0.0255) loss_box_reg: 0.0332
(0.0402) loss_objectness: 0.0014 (0.0020) loss_rpn_box_reg: 0.0030
(0.0068) time: 0.3708 data: 0.0632 max mem: 2795
Epoch: [16] [120/170] eta: 0:00:18 lr: 0.000000 loss: 0.0677
(0.0739) loss_classifier: 0.0212 (0.0252) loss_box_reg: 0.0369
(0.0403) loss_objectness: 0.0002 (0.0019) loss_rpn_box_reg: 0.0025
(0.0064) time: 0.3688 data: 0.0653 max mem: 2795
Epoch: [16] [140/170] eta: 0:00:10 lr: 0.000000 loss: 0.0547
(0.0725) loss_classifier: 0.0217 (0.0249) loss_box_reg: 0.0302
(0.0395) loss_objectness: 0.0002 (0.0020) loss_rpn_box_reg: 0.0013
(0.0061) time: 0.3613 data: 0.0608 max mem: 2795
Epoch: [16] [160/170] eta: 0:00:03 lr: 0.000000 loss: 0.0603
(0.0727) loss_classifier: 0.0234 (0.0251) loss_box_reg: 0.0361
(0.0395) loss_objectness: 0.0002 (0.0019) loss_rpn_box_reg: 0.0020
(0.0061) time: 0.3561 data: 0.0600 max mem: 2795
Epoch: [16] [169/170] eta: 0:00:00 lr: 0.000000 loss: 0.0551
(0.0725) loss_classifier: 0.0234 (0.0250) loss_box_reg: 0.0315
(0.0395) loss_objectness: 0.0004 (0.0019) loss_rpn_box_reg: 0.0022
(0.0061) time: 0.3587 data: 0.0623 max mem: 2795
Epoch: [16] Total time: 0:01:01 (0.3639 s / it)
mAP:0.467
  AP at IoU level [0.50]: 0.714
  AP at IoU level [0.55]: 0.697
  AP at IoU level [0.60]: 0.676
  AP at IoU level [0.65]: 0.668
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.574
  AP at IoU level [0.80]: 0.422
  AP at IoU level [0.85]: 0.250
  AP at IoU level [0.90]: 0.040
  AP at IoU level [0.95]: 0.000
Epoch: [17] [ 0/170] eta: 0:01:04 lr: 0.000000 loss: 0.1328
(0.1328) loss_classifier: 0.0380 (0.0380) loss_box_reg: 0.0741
(0.0741) loss_objectness: 0.0057 (0.0057) loss_rpn_box_reg: 0.0150
(0.0150) time: 0.3776 data: 0.0643 max mem: 2795
Epoch: [17] [ 20/170] eta: 0:00:55 lr: 0.000000 loss: 0.0692
(0.0774) loss_classifier: 0.0247 (0.0251) loss_box_reg: 0.0408
(0.0432) loss_objectness: 0.0006 (0.0023) loss_rpn_box_reg: 0.0039
```

```
(0.0068) time: 0.3721 data: 0.0653 max mem: 2795
Epoch: [17] [ 40/170] eta: 0:00:47 lr: 0.000000 loss: 0.0638
(0.0821) loss_classifier: 0.0257 (0.0266) loss_box_reg: 0.0300
(0.0453) loss_objectness: 0.0011 (0.0026) loss_rpn_box_reg: 0.0027
(0.0077) time: 0.3589 data: 0.0649 max mem: 2795
Epoch: [17] [ 60/170] eta: 0:00:40 lr: 0.000000 loss: 0.0686
(0.0797) loss_classifier: 0.0212 (0.0264) loss_box_reg: 0.0367
(0.0436) loss_objectness: 0.0009 (0.0025) loss_rpn_box_reg: 0.0021
(0.0072) time: 0.3622 data: 0.0571 max mem: 2795
Epoch: [17] [ 80/170] eta: 0:00:32 lr: 0.000000 loss: 0.0609
(0.0795) loss_classifier: 0.0242 (0.0267) loss_box_reg: 0.0335
(0.0433) loss_objectness: 0.0007 (0.0026) loss_rpn_box_reg: 0.0018
(0.0068) time: 0.3687 data: 0.0641 max mem: 2795
Epoch: [17] [100/170] eta: 0:00:25 lr: 0.000000 loss: 0.0459
(0.0754) loss_classifier: 0.0207 (0.0258) loss_box_reg: 0.0224
(0.0407) loss_objectness: 0.0001 (0.0023) loss_rpn_box_reg: 0.0008
(0.0067) time: 0.3521 data: 0.0597 max mem: 2795
Epoch: [17] [120/170] eta: 0:00:18 lr: 0.000000 loss: 0.0442
(0.0723) loss_classifier: 0.0202 (0.0250) loss_box_reg: 0.0207
(0.0391) loss_objectness: 0.0001 (0.0021) loss_rpn_box_reg: 0.0009
(0.0061) time: 0.3604 data: 0.0587 max mem: 2795
Epoch: [17] [140/170] eta: 0:00:10 lr: 0.000000 loss: 0.0499
(0.0700) loss_classifier: 0.0173 (0.0245) loss_box_reg: 0.0283
(0.0380) loss_objectness: 0.0001 (0.0019) loss_rpn_box_reg: 0.0018
(0.0056) time: 0.3609 data: 0.0568 max mem: 2795
Epoch: [17] [160/170] eta: 0:00:03 lr: 0.000000 loss: 0.0739
(0.0717) loss_classifier: 0.0252 (0.0248) loss_box_reg: 0.0441
(0.0390) loss_objectness: 0.0006 (0.0020) loss_rpn_box_reg: 0.0024
(0.0060) time: 0.3563 data: 0.0614 max mem: 2795
Epoch: [17] [169/170] eta: 0:00:00 lr: 0.000000 loss: 0.0739
(0.0719) loss_classifier: 0.0252 (0.0248) loss_box_reg: 0.0407
(0.0392) loss_objectness: 0.0006 (0.0020) loss_rpn_box_reg: 0.0024
(0.0060) time: 0.3572 data: 0.0606 max mem: 2795
Epoch: [17] Total time: 0:01:01 (0.3611 s / it)
mAP:0.464
  AP at IoU level [0.50]: 0.707
  AP at IoU level [0.55]: 0.697
  AP at IoU level [0.60]: 0.677
  AP at IoU level [0.65]: 0.669
  AP at IoU level [0.70]: 0.632
  AP at IoU level [0.75]: 0.575
  AP at IoU level [0.80]: 0.402
  AP at IoU level [0.85]: 0.244
  AP at IoU level [0.90]: 0.040
  AP at IoU level [0.95]: 0.000
Epoch: [18] [ 0/170] eta: 0:00:55 lr: 0.000000 loss: 0.0357
(0.0357) loss_classifier: 0.0176 (0.0176) loss_box_reg: 0.0176
(0.0176) loss_objectness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0004
(0.0004) time: 0.3290 data: 0.0692 max mem: 2795
Epoch: [18] [ 20/170] eta: 0:00:54 lr: 0.000000 loss: 0.0599
```

```
(0.0769) loss_classifier: 0.0234 (0.0239) loss_box_reg: 0.0361
(0.0439) loss_objectness: 0.0004 (0.0016) loss_rpn_box_reg: 0.0022
(0.0075) time: 0.3619 data: 0.0612 max mem: 2795
Epoch: [18] [ 40/170] eta: 0:00:47 lr: 0.000000 loss: 0.0690
(0.0832) loss_classifier: 0.0217 (0.0256) loss_box_reg: 0.0399
(0.0475) loss_objectness: 0.0005 (0.0026) loss_rpn_box_reg: 0.0022
(0.0074) time: 0.3691 data: 0.0630 max mem: 2795
Epoch: [18] [ 60/170] eta: 0:00:39 lr: 0.000000 loss: 0.0560
(0.0752) loss_classifier: 0.0228 (0.0244) loss_box_reg: 0.0276
(0.0425) loss_objectness: 0.0002 (0.0022) loss_rpn_box_reg: 0.0014
(0.0061) time: 0.3593 data: 0.0599 max mem: 2795
Epoch: [18] [ 80/170] eta: 0:00:32 lr: 0.000000 loss: 0.0458
(0.0700) loss_classifier: 0.0199 (0.0239) loss_box_reg: 0.0248
(0.0389) loss_objectness: 0.0001 (0.0020) loss_rpn_box_reg: 0.0008
(0.0052) time: 0.3637 data: 0.0594 max mem: 2795
Epoch: [18] [100/170] eta: 0:00:25 lr: 0.000000 loss: 0.0484
(0.0689) loss_classifier: 0.0209 (0.0239) loss_box_reg: 0.0266
(0.0375) loss_objectness: 0.0002 (0.0018) loss_rpn_box_reg: 0.0012
(0.0058) time: 0.3656 data: 0.0555 max mem: 2795
Epoch: [18] [120/170] eta: 0:00:18 lr: 0.000000 loss: 0.0553
(0.0694) loss_classifier: 0.0200 (0.0239) loss_box_reg: 0.0341
(0.0380) loss_objectness: 0.0003 (0.0019) loss_rpn_box_reg: 0.0030
(0.0056) time: 0.3736 data: 0.0637 max mem: 2795
Epoch: [18] [140/170] eta: 0:00:10 lr: 0.000000 loss: 0.0593
(0.0691) loss_classifier: 0.0246 (0.0243) loss_box_reg: 0.0345
(0.0378) loss_objectness: 0.0002 (0.0017) loss_rpn_box_reg: 0.0018
(0.0052) time: 0.3643 data: 0.0627 max mem: 2795
Epoch: [18] [160/170] eta: 0:00:03 lr: 0.000000 loss: 0.0658
(0.0713) loss_classifier: 0.0254 (0.0247) loss_box_reg: 0.0314
(0.0390) loss_objectness: 0.0009 (0.0018) loss_rpn_box_reg: 0.0026
(0.0058) time: 0.3610 data: 0.0598 max mem: 2795
Epoch: [18] [169/170] eta: 0:00:00 lr: 0.000000 loss: 0.0602
(0.0712) loss_classifier: 0.0248 (0.0247) loss_box_reg: 0.0298
(0.0388) loss_objectness: 0.0001 (0.0018) loss_rpn_box_reg: 0.0016
(0.0059) time: 0.3602 data: 0.0611 max mem: 2795
Epoch: [18] Total time: 0:01:01 (0.3644 s / it)
mAP:0.466
  AP at IoU level [0.50]: 0.707
  AP at IoU level [0.55]: 0.698
  AP at IoU level [0.60]: 0.676
  AP at IoU level [0.65]: 0.668
  AP at IoU level [0.70]: 0.631
  AP at IoU level [0.75]: 0.574
  AP at IoU level [0.80]: 0.420
  AP at IoU level [0.85]: 0.245
  AP at IoU level [0.90]: 0.040
  AP at IoU level [0.95]: 0.000
Epoch: [19] [ 0/170] eta: 0:01:03 lr: 0.000000 loss: 0.0642
(0.0642) loss_classifier: 0.0244 (0.0244) loss_box_reg: 0.0336
(0.0336) loss_objectness: 0.0022 (0.0022) loss_rpn_box_reg: 0.0040
```

(0.0040) time: 0.3710 data: 0.0608 max mem: 2795
Epoch: [19] [20/170] eta: 0:00:53 lr: 0.000000 loss: 0.0741
(0.0830) loss_classifier: 0.0262 (0.0284) loss_box_reg: 0.0404
(0.0445) loss_objectness: 0.0006 (0.0034) loss_rpn_box_reg: 0.0031
(0.0067) time: 0.3566 data: 0.0596 max mem: 2795
Epoch: [19] [40/170] eta: 0:00:47 lr: 0.000000 loss: 0.0513
(0.0805) loss_classifier: 0.0202 (0.0275) loss_box_reg: 0.0290
(0.0439) loss_objectness: 0.0002 (0.0029) loss_rpn_box_reg: 0.0009
(0.0063) time: 0.3673 data: 0.0636 max mem: 2795
Epoch: [19] [60/170] eta: 0:00:39 lr: 0.000000 loss: 0.0682
(0.0789) loss_classifier: 0.0240 (0.0271) loss_box_reg: 0.0399
(0.0429) loss_objectness: 0.0010 (0.0026) loss_rpn_box_reg: 0.0020
(0.0064) time: 0.3660 data: 0.0651 max mem: 2795
Epoch: [19] [80/170] eta: 0:00:32 lr: 0.000000 loss: 0.0649
(0.0775) loss_classifier: 0.0239 (0.0265) loss_box_reg: 0.0369
(0.0420) loss_objectness: 0.0005 (0.0025) loss_rpn_box_reg: 0.0014
(0.0066) time: 0.3610 data: 0.0611 max mem: 2795
Epoch: [19] [100/170] eta: 0:00:25 lr: 0.000000 loss: 0.0563
(0.0779) loss_classifier: 0.0235 (0.0266) loss_box_reg: 0.0307
(0.0419) loss_objectness: 0.0003 (0.0025) loss_rpn_box_reg: 0.0022
(0.0069) time: 0.3684 data: 0.0581 max mem: 2795
Epoch: [19] [120/170] eta: 0:00:18 lr: 0.000000 loss: 0.0449
(0.0758) loss_classifier: 0.0210 (0.0259) loss_box_reg: 0.0270
(0.0407) loss_objectness: 0.0008 (0.0025) loss_rpn_box_reg: 0.0011
(0.0066) time: 0.3596 data: 0.0656 max mem: 2795
Epoch: [19] [140/170] eta: 0:00:10 lr: 0.000000 loss: 0.0500
(0.0729) loss_classifier: 0.0185 (0.0252) loss_box_reg: 0.0276
(0.0393) loss_objectness: 0.0000 (0.0023) loss_rpn_box_reg: 0.0008
(0.0060) time: 0.3646 data: 0.0571 max mem: 2795
Epoch: [19] [160/170] eta: 0:00:03 lr: 0.000000 loss: 0.0542
(0.0725) loss_classifier: 0.0211 (0.0252) loss_box_reg: 0.0287
(0.0392) loss_objectness: 0.0005 (0.0022) loss_rpn_box_reg: 0.0015
(0.0059) time: 0.3575 data: 0.0616 max mem: 2795
Epoch: [19] [169/170] eta: 0:00:00 lr: 0.000000 loss: 0.0438
(0.0719) loss_classifier: 0.0181 (0.0250) loss_box_reg: 0.0253
(0.0388) loss_objectness: 0.0002 (0.0022) loss_rpn_box_reg: 0.0013
(0.0059) time: 0.3611 data: 0.0597 max mem: 2795
Epoch: [19] Total time: 0:01:01 (0.3631 s / it)
mAP:0.468

AP at IoU level [0.50]: 0.715
AP at IoU level [0.55]: 0.698
AP at IoU level [0.60]: 0.677
AP at IoU level [0.65]: 0.669
AP at IoU level [0.70]: 0.632
AP at IoU level [0.75]: 0.575
AP at IoU level [0.80]: 0.420
AP at IoU level [0.85]: 0.245
AP at IoU level [0.90]: 0.046
AP at IoU level [0.95]: 0.000

Saving the Model

```
from datetime import datetime
# Save model with current date
now = datetime.now()
d = now.strftime("%Y_%b_%d_%Hh_%mm")
PATH = 'model_'+d+'.pt'

torch.save(model.state_dict(), PATH)
```

Evaluate and Predict on Test Set

```
# Get saved model
model_eval = model.load_state_dict(torch.load(PATH))
```

Evaluation:

```
# put the model in evaluation mode
model.eval()

# Evaluate the model
evaluate(model, loader_test, device=device)
```

mAP:0.419

```
AP at IoU level [0.50]: 0.645
AP at IoU level [0.55]: 0.639
AP at IoU level [0.60]: 0.629
AP at IoU level [0.65]: 0.611
AP at IoU level [0.70]: 0.566
AP at IoU level [0.75]: 0.485
AP at IoU level [0.80]: 0.370
AP at IoU level [0.85]: 0.196
AP at IoU level [0.90]: 0.044
AP at IoU level [0.95]: 0.002
```

```
(0.41878714099522474,
 [0.6450177478928426,
  0.6389704014059845,
  0.629120241522798,
  0.6111861989017132,
  0.5657634863451461,
  0.48528152182043605,
  0.37007749023925945,
  0.19647219081539297,
  0.044251811806310576,
  0.0017303192023640987])
```

Test prediction on random image.

```
# Make prediction on random image
n = randint(0, dataset_test.len)
img, target = dataset_test[n]
with torch.no_grad():
```

```


prediction = model([img.to(device)])[0]

# Non max suppression to reduce the number of bounding boxes
nms_prediction = apply_nms(prediction, iou_thresh=0.5)
# Remove low score boxes below score_thresh
filtered_prediction = remove_low_score_bb(nms_prediction,
score_thresh=0.3)

# Draw bounding boxes
draw_bounding_boxes(img.detach().cpu(), target=target,
prediction=filtered_prediction)

```

400x208



Evaluation from coco tools

```

from engine import evaluate as eval
eval(model, loader_test, device=device)

creating index...
index created!
Test: [ 0/38] eta: 0:00:07 model_time: 0.1497 (0.1497)
evaluator_time: 0.0055 (0.0055) time: 0.2033 data: 0.0461 max mem:
2795
Test: [37/38] eta: 0:00:00 model_time: 0.1278 (0.1290)
evaluator_time: 0.0088 (0.0208) time: 0.2180 data: 0.0592 max mem:
2795
Test: Total time: 0:00:08 (0.2127 s / it)
Averaged stats: model_time: 0.1278 (0.1290) evaluator_time: 0.0088
(0.0208)
Accumulating evaluation results...
DONE (t=0.07s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |

```

```

maxDets=100 ] = 0.470
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.757
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.519
Average Precision (AP) @[ IoU=0.50:0.95 | area=  small |
maxDets=100 ] = 0.347
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.694
Average Precision (AP) @[ IoU=0.50:0.95 | area=  large |
maxDets=100 ] = 0.799
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.277
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.518
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.540
Average Recall    (AR) @[ IoU=0.50:0.95 | area=  small |
maxDets=100 ] = 0.422
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.760
Average Recall    (AR) @[ IoU=0.50:0.95 | area=  large |
maxDets=100 ] = 0.823

```

```
<coco_eval.CocoEvaluator at 0x7fc2084b6e90>
```

VGG19 Approach

How to create a Deep Learning face mask classifier for COVID-19 in public spaces

Introduction

The [CDC](#) continues to monitor the spread of COVID-19 and advises people who are completely vaccinated as well as those who are not fully vaccinated to wear face masks. When visiting the doctor's office, hospitals, or long-term care institutions, the CDC recommends wearing masks and keeping a safe distance.

Manually monitoring people entering such institutions is tedious and requires workforce. In this tutorial, we will learn how we can automate this process through deep learning techniques which will automatically detect people not wearing masks to prevent their entry.

Creating the mask detection deep learning model

We will now look into building a Deep Learning model to predict (detect) if a person is violating the rules by not wearing a mask in public spaces.

Step 1: Importing the necessary Python libraries

```
import numpy as np # linear algebra
import cv2 # opencv
import matplotlib.pyplot as plt # image plotting
# keras
from keras import Sequential
from keras.layers import Flatten, Dense
from keras.applications.vgg19 import VGG19
from keras.applications.vgg19 import preprocess_input
from keras.preprocessing.image import ImageDataGenerator
```

Step 2: Getting the data

For the training data, we are using the face mask detection data from [here](#). The dataset contains 12 thousand images divided into Test, Train, and Validation sets which were scraped from Google and the CelebFace dataset created by [Jessica Li](#).

```
# Load train and test set
train_dir = "/content/FaceMaskDetection12k/Train"
test_dir = "/content/FaceMaskDetection12k/Test"
val_dir = "/content/FaceMaskDetection12k/Validation"
```

Step 3: Reading a sample image and performing face detection

We will now read in a sample image from a busy airport and perform face detection using haar cascade classifier. The [Haar cascade classifier](#), originally known as the Viola-Jones Face Detection Technique is a object detection algorithm for detecting faces in images or real-time video.

Viola and Jones proposed edge or line detection features in their research paper "Rapid Object Detection using a Boosted Cascade of Simple Features," published in 2001. The algorithm is given a large number of positive photos with faces and a large number of negative images with no faces. The model developed as a result of this training can be found in the OpenCV GitHub [repository](#).

```
# Read a sample image
img =
cv2.imread("../input/face-mask-detection/images/maksssksksss352.png")

# Keep a copy of coloured image
orig_img = cv2.cvtColor(img, cv2.COLOR_RGB2BGR) # colored output
image

# Convert image to grayscale
img = cv2.cvtColor(img, cv2.IMREAD_GRAYSCALE)

# loading haarcascade_frontalface_default.xml
face_detection_model = cv2.CascadeClassifier("../input/haar-cascades-
for-face-detection/haarcascade_frontalface_default.xml")
```

```
# detect faces in the given image
return_faces = face_detection_model.detectMultiScale(
    img, scaleFactor=1.08, minNeighbors=4
) # returns a list of (x,y,w,h) tuples

# plotting the returned values
for (x, y, w, h) in return_faces:
    cv2.rectangle(orig_img, (x, y), (x + w, y + h), (0, 0, 255), 1)

plt.figure(figsize=(12, 12))
plt.imshow(orig_img) # display the image

<matplotlib.image.AxesImage at 0x7fb4fa097450>
```



Step 4: Data preprocessing for building the mask detection Keras model

We will now pass our datasets into Keras `ImageDataGenerator()` to perform some preliminary data augmentation steps such as rescaling.

```
# Data preprocessing
# Train data
datagenerator = ImageDataGenerator(
    rescale=1.0 / 255, horizontal_flip=True, zoom_range=0.2,
    shear_range=0.2
)
train_generator = datagenerator.flow_from_directory(
    directory=train_dir, target_size=(128, 128),
    class_mode="categorical", batch_size=32
)

# Validation data
val_generator = datagenerator.flow_from_directory(
    directory=val_dir, target_size=(128, 128),
    class_mode="categorical", batch_size=32
)

# Test data
test_generator = datagenerator.flow_from_directory(
    directory=val_dir, target_size=(128, 128),
    class_mode="categorical", batch_size=32
)
```

```
Found 10000 images belonging to 2 classes.
Found 800 images belonging to 2 classes.
Found 800 images belonging to 2 classes.
```

Step 5: Create the mask detection transfer learning model using Keras

We are building the deep learning classifier using the VGG19 transfer learning model. The VGG19 model is the successor of AlexNet, a variation of the VGG model named after the group named as Visual Geometry Group at Oxford which created it. It is a deep CNN consisting of 19 layers (16 convolution layers, 3 Fully connected layer, 5 MaxPool layers and 1 SoftMax layer) used to classify images.

It has been trained on [ImageNet](#), a picture database with 14,197,122 images structured according to the WordNet hierarchy.

VGG19 Architecture

```
# Initializing the VGG19 model
vgg19_model = VGG19(weights="imagenet", include_top=False,
input_shape=(128, 128, 3))

for layer in vgg19_model.layers:
    layer.trainable = False
```

```
# Initialize a sequential model
```

```
model = Sequential()
model.add(vgg19_model)
model.add(Flatten())
model.add(Dense(2, activation="sigmoid"))
model.summary()
```

```
# Compiling the model
```

```
model.compile(optimizer="adam", loss="categorical_crossentropy",
metrics="accuracy")
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80142336/80134624 [=====] - 1s 0us/step
80150528/80134624 [=====] - 1s 0us/step
Model: "sequential"
```

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 4, 4, 512)	20024384
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 2)	16386

```
=====
Total params: 20,040,770
Trainable params: 16,386
Non-trainable params: 20,024,384
=====
```

Step 6: Train the model

We will now train our neural network model for 20 epochs.

```
# Fit the model on train data along with validation data
```

```
model_history = model.fit_generator(
    generator=train_generator,
    steps_per_epoch=len(train_generator) // 32,
    epochs=20,
    validation_data=val_generator,
    validation_steps=len(val_generator) // 32,
)
```

```
/opt/conda/lib/python3.7/site-packages/keras/engine/training.py:1972:
UserWarning: `Model.fit_generator` is deprecated and will be removed
in a future version. Please use `Model.fit`, which supports
generators.
```


Epoch 1/20

14:36:57.564987: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005

9/9 [=====] - 11s 370ms/step - loss: 0.5636 -
accuracy: 0.7465

Epoch 2/20

9/9 [=====] - 3s 348ms/step - loss: 0.2521 -
accuracy: 0.9167

Epoch 3/20

9/9 [=====] - 3s 389ms/step - loss: 0.1610 -
accuracy: 0.9444

Epoch 4/20

9/9 [=====] - 3s 344ms/step - loss: 0.1345 -
accuracy: 0.9514

Epoch 5/20

9/9 [=====] - 3s 324ms/step - loss: 0.1158 -
accuracy: 0.9688

Epoch 6/20

9/9 [=====] - 3s 348ms/step - loss: 0.1159 -
accuracy: 0.9618

Epoch 7/20

9/9 [=====] - 3s 326ms/step - loss: 0.1098 -
accuracy: 0.9722

Epoch 8/20

9/9 [=====] - 3s 353ms/step - loss: 0.1017 -
accuracy: 0.9688

Epoch 9/20

9/9 [=====] - 3s 311ms/step - loss: 0.0774 -
accuracy: 0.9757

Epoch 10/20

9/9 [=====] - 3s 314ms/step - loss: 0.0698 -
accuracy: 0.9826

Epoch 11/20

9/9 [=====] - 3s 370ms/step - loss: 0.0733 -
accuracy: 0.9792

Epoch 12/20

9/9 [=====] - 3s 285ms/step - loss: 0.0676 -
accuracy: 0.9757

Epoch 13/20

9/9 [=====] - 3s 293ms/step - loss: 0.0744 -
accuracy: 0.9792

Epoch 14/20

9/9 [=====] - 3s 313ms/step - loss: 0.0690 -
accuracy: 0.9826

Epoch 15/20

9/9 [=====] - 3s 280ms/step - loss: 0.0955 -
accuracy: 0.9618

Epoch 16/20

9/9 [=====] - 3s 295ms/step - loss: 0.0488 -

```

accuracy: 0.9861
Epoch 17/20
9/9 [=====] - 3s 280ms/step - loss: 0.0519 -
accuracy: 0.9826
Epoch 18/20
9/9 [=====] - 3s 284ms/step - loss: 0.0456 -
accuracy: 0.9931
Epoch 19/20
9/9 [=====] - 2s 268ms/step - loss: 0.0581 -
accuracy: 0.9792
Epoch 20/20
9/9 [=====] - 3s 284ms/step - loss: 0.0607 -
accuracy: 0.9722

```

Step 7: Evaluate the model performance on test set

Evaluate model performance on test data

```

model_loss, model_acc = model.evaluate(test_generator)
print("Model has a loss of %.2f and accuracy %.2f%%" % (model_loss,
model_acc*100))

```

```

25/25 [=====] - 9s 350ms/step - loss: 0.0678
- accuracy: 0.9775
Model has a loss of 0.07 and accuracy 97.75%

```

Step 8: Save the model

We can also choose to save the trained model as a h5 file for future use.

```

model.save('data/saved_model.h5')

```

Step 9: Test the model on the sample image

We will now test the trained model on our use case for detecting faces and masks for a group of people. We take the detected face crops of the faces detected in the image and then predict the mask or no mask using the model trained.

label for mask detection

```

mask_det_label = {0: "Mask", 1: "No Mask"}
mask_det_label_colour = {0: (0, 255, 0), 1: (255, 0, 0)}
pad_y = 1 # padding for result text

```

```

main_img = cv2.cvtColor(img, cv2.COLOR_RGB2BGR) # colored output image

```

For detected faces in the image

```

for i in range(len(return_faces)):
    (x, y, w, h) = return_faces[i]
    cropped_face = main_img[y : y + h, x : x + w]
    cropped_face = cv2.resize(cropped_face, (128, 128))
    cropped_face = np.reshape(cropped_face, [1, 128, 128, 3]) / 255.0
    mask_result = model.predict(cropped_face) # make model prediction

```

```

    print_label = mask_det_label[mask_result.argmax()] # get mask/no
mask based on prediction
    label_colour = mask_det_label_colour[mask_result.argmax()] # green
for mask, red for no mask

    # Print result
    (t_w, t_h), _ = cv2.getTextSize(
        print_label, cv2.FONT_HERSHEY_SIMPLEX, 0.4, 1
    ) # getting the text size


    cv2.rectangle(
        main_img,
        (x, y + pad_y),
        (x + t_w, y - t_h - pad_y - 6),
        label_colour,
        -1,
    ) # draw rectangle

    cv2.putText(
        main_img,
        print_label,
        (x, y - 6),
        cv2.FONT_HERSHEY_DUPLEX,
        0.4,
        (255, 255, 255), # white
        1,
    ) # print text

    cv2.rectangle(
        main_img,
        (x, y),
        (x + w, y + h),
        label_colour,
        1,
    ) # draw bounding box on face

plt.figure(figsize=(10, 10))
plt.imshow(main_img) # display image
<matplotlib.image.AxesImage at 0x7fb4b8b52910>

```



We can see that the model is correctly detecting faces and classifying them as mask and no mask.